

ARE THERE FACTIVE PREDICATES? AN EMPIRICAL INVESTIGATION

JUDITH DEGEN

Stanford University

JUDITH TONHAUSER

University of Stuttgart

Properties of the content of the clausal complement have long been assumed to distinguish factive predicates like *know* from nonfactive ones like *think* (Kiparsky & Kiparsky 1970, *inter alia*). There is, however, disagreement about which properties define factive predicates, as well as uncertainty about whether the content of the complement of particular predicates exhibits the properties attributed to the content of the complement of factive predicates. This has led to a lack of consensus about which predicates are factive, a troublesome situation given the central role that factivity plays in linguistic theorizing. This article reports six experiments designed to investigate two critical properties of the content of the complement of clause-embedding predicates, namely projection and entailment, with the goal of establishing whether these properties identify a class of factive predicates. We find that factive predicates are more heterogeneous than previously assumed and that there is little empirical support from these experiments for the assumed categorical distinction between factive and nonfactive predicates. We discuss implications of our results for formal analyses of presuppositions, one area where factivity has played a central role. We propose that projection is sensitive to meaning distinctions between clause-embedding predicates that are more fine-grained than factivity.*

Keywords: factivity, factive predicates, presupposition, projection, entailment, experiments

1. INTRODUCTION. Clause-embedding predicates are predicates like *know* and *think* that take a finite clause as an argument, as shown in 1, where the finite clause (*that*) *it is raining* realizes the direct object. Kiparsky & Kiparsky 1970 proposed that there is a class of clause-embedding predicates called FACTIVE predicates, like *know*, that distinguish themselves from other predicates, like *think*, in that the content of their clausal complement is presupposed, that is, ‘[t]he speaker presupposes that the embedded clause expresses a true proposition’ (p. 147). Thus, a speaker who utters 1 with *know* is typically taken to presuppose that it is raining, in contrast to a speaker who utters 1 with *think*.

(1) Sam {knows/thinks} that it is raining.

To diagnose whether a predicate is factive, that is, whether the content of its complement (henceforth CC) is presupposed, Kiparsky and Kiparsky (1970) relied on the projection diagnostic (see also e.g. Langendoen & Savin 1971). Specifically, they assumed that the CC of a clause-embedding predicate is presupposed if the speaker is typically committed to the truth of the CC not only when uttering an affirmative matrix sentence like 1 but also when uttering negated or polar-question variants of that sentence, like those in 2a and 2b, respectively. Thus, *know* was categorized in Kiparsky & Kiparsky 1970 as a factive predicate because a speaker typically is committed to the CC not only when uttering the affirmative sentence in 1 with *know*, but also when uttering its negated variant in 2a or the polar-question variant in 2b. By contrast, *think* was categorized as nonfactive because the CC is not presupposed; that is, a speaker who utters the variants in 1 and 2a–b with *think* is not typically committed to the truth of the CC.¹

* This research has benefited tremendously from sustained interaction with David Beaver, Mandy Simons, and Craige Roberts. For helpful comments we also thank audiences at the 2018 meeting of XPRAG.de in Graz, the 20th Annual SemFest at Stanford, SALT 30 at MIT, and at the universities in Tübingen, Saarbrücken, Harvard, British Columbia, and Bochum. We gratefully acknowledge financial support for this research from National Science Foundation grant BCS-1452674 (JT).

¹ Factive predicates were also taken to be distinguished by syntactic properties in Kiparsky & Kiparsky 1970. We ignore syntax here because Kiparsky and Kiparsky (1970, n. 3) already pointed out that syntactic and semantic factivity are not perfectly aligned: *know* and *realize*, for instance, were taken to be syntactically

- (2) a. Sam doesn't {know/think} that it is raining.
 b. Does Sam {know/think} that it is raining?

Today, some fifty years after Kiparsky & Kiparsky 1970, properties of the CC are still taken to categorically distinguish factive predicates from others, in English and other languages. However, as we show below, there is disagreement about which properties define factive predicates and uncertainty about whether the CCs of particular predicates exhibit the properties attributed to factive predicates. As a result, there is no consensus on which predicates are factive. This lack of consensus is problematic given the central role that the categorization plays in linguistic theorizing on topics as diverse as projection (e.g. Karttunen & Peters 1979, van der Sandt 1992), embedded questions (e.g. Hintikka 1975, Guerzoni & Sharvit 2007, Spector & Egré 2015), extraction and island effects (e.g. Hukari & Levine 1995, Rooryck 2000, Abrusán 2014), exclamatives (e.g. Zanuttini & Portner 2003), control (e.g. Landau 2001), mood (e.g. Givón 1995, van Gelderen 2004, Heycock 2006, Giannakidou & Mari 2015), negative polarity and free choice (e.g. Giannakidou 1998, 2001), predicates of personal taste (e.g. Lasersohn 2009), and language acquisition and theory of mind (e.g. de Villiers 2005). This article reports six experiments designed to investigate critical properties of the CC of clause-embedding predicates, with the goal of investigating whether these properties identify a coherent class of factive predicates. What we find, in a nutshell, is that predicates canonically deemed factive are more heterogeneous than previously assumed and that there is little empirical support from these experiments for the assumed categorical distinction between factive and nonfactive predicates.

The remainder of this introduction shows that there is disagreement about which properties define factive predicates (§1.1) and uncertainty about whether the CCs of particular predicates exhibit the properties attributed to the CCs of factive predicates (§1.2).

1.1. DISAGREEMENT ABOUT THE DEFINITION OF FACTIVE PREDICATES. One reason for the lack of consensus about which predicates are factive is disagreement about their definition. As mentioned above, Kiparsky & Kiparsky 1970 defined factive predicates as those for which the CC is presupposed (see also e.g. Karttunen 1971a,b), with projection diagnosing presupposed CCs. Thus, on this definition, a predicate is factive if and only if the speaker is typically committed to the truth of the CC not only in affirmative matrix sentences, but also in negated and polar-question variants. More recent work, however, assumes that the CC of a factive predicate is not only presupposed but also entailed (e.g. Gazdar 1979a:119–23, Schlenker 2010:139, Abrusán 2011, Anand & Hacquard 2014:71, Spector & Egré 2015, n. 7).² In other words, for a predicate to be factive, not only does the speaker have to presuppose the CC, but the truth of the CC also has to be *indefeasibly inferred* from affirmative matrix sentences: that is, the truth

nonfactive but semantically factive. For recent discussion see White & Rawlins 2018, Djärv 2019, and references therein.

² For instance, according to Beaver (2001:66f.), Gazdar (1979a:119–23) ‘describes the inferences associated with factive verbs … as being *indefeasible* in simple affirmative sentences’, that is, ‘entailments’. And Schlenker (2010:139), in discussing CCs, pointed to ‘the pattern of inference which is characteristic of presuppositions: an entailment of the positive sentence is preserved under negation and in questions’. Authors are not always clear about which definition they subscribe to. For instance, Chierchia and McConnell-Ginet (1990:355) used *realize* to illustrate that ‘[a] sentence can both entail and presuppose another sentence’: *Joan realizes that syntax deals with sentence structure ‘both entails and presupposes’* (*ibid.*) that syntax deals with sentence structure. It is not clear, however, whether they assumed this pattern for factive predicates more generally or whether they assumed that a predicate is factive as long as its CC is presupposed.

of the CC has to follow for any competent language user, not just the speaker, from a true affirmative matrix sentence in which the CC is embedded under the predicate.³ The two definitions are given in 3.

- (3) Two definitions of factive predicates: A clause-embedding predicate is factive if and only if the content of its clausal complement is
 - a. presupposed. (e.g. Kiparsky & Kiparsky 1970, Karttunen 1971a,b)
 - b. presupposed and entailed. (e.g. Gazdar 1979a, Schlenker 2010, Abrusán 2011, Anand & Hacquard 2014, Spector & Egré 2015)

These two definitions come apart on some predicates, including *take into account* and *make clear*, shown in 4 and 5. These predicates were considered factive in Kiparsky & Kiparsky 1970, on definition 3a, because the CC projects: that is, the speaker is committed to the truth of the CC, that the groom is unreliable, not just when uttering one of the affirmative matrix sentences in 4a, but also when uttering one of the polar-question variants in 4b. But true affirmative matrix sentences with these predicates do not entail the CC, as shown by the acceptability of 5. In this example, the truth of the CC is not an indefeasible inference of the affirmative matrix sentence with one of these predicates because the affirmative matrix sentence can be followed up, without contradiction, by *because the groom is a very reliable person*. If the CC was entailed by the affirmative matrix sentence, 5 should be contradictory. Thus, *take into account* and *make clear*, which were taken to be factive on definition 3a, are not factive on definition 3b.

- (4) a. Joan {took into account/made clear} that the groom is unreliable.
- b. Did Joan {take into account/make clear} that the groom is unreliable?
- (5) Joan {took into account/made clear} that the groom is unreliable, but she was wrong to do so, because the groom is a very reliable person.

There are more predicates on which the two definitions come apart, as shown in §1.2, which illustrates disagreements about whether the CC of particular clause-embedding predicates is presupposed and entailed.

1.2. PROPERTIES OF THE CONTENT OF THE COMPLEMENT OF CLAUSE-EMBEDDING PREDICATES.

PRESUPPOSED. The two definitions in 3 share the assumption that the CC of factive predicates is presupposed. In addition to being shared by the two definitions, this assumption is also shared by different formal analyses of presuppositions: that is, the CC of factive predicates is assumed to be presupposed regardless of whether presuppositions are formally analyzed as conventionally coded constraints on the common ground (e.g. Heim 1983, van der Sandt 1992), as derived from lexical alternatives (e.g. Abusch 2010, Romoli 2015), or as pragmatically derived (e.g. Abrusán 2011, Simons et al. 2017). In

³ A content *c* is an entailment of sentence *S* if and only if *c* is the case whenever *S* is true, that is, if and only if *c* is indefeasibly implied by *S*. For instance, the content that Sam has a cat is an entailment of the sentence *Sam has a white cat*. A hallmark of presupposed content is projection: a content *c* is a presupposition if *c* is typically implied by an utterance of *S* as well as by utterances of so-called family-of-sentences variants of *S*, including the polar-question variant or negation of *S* (see e.g. Chierchia & McConnell-Ginet 1990:§3.1). For instance, the content that Sam has a parrot is typically implied not just by an utterance of the sentence *Sam's parrot is loud*, but also by utterances of the polar-question variant *Is Sam's parrot loud?* and the negated variant *Sam's parrot isn't loud*. (See §1.2 for how to distinguish presuppositions from conventional implicatures, another type of projective content.) From the definitions of presupposed and entailed content, and from the ways in which projection and entailment are diagnosed, it follows that presupposed content may, but need not, be entailed, and entailed content may, but need not, be presupposed. Thus, content that is implied by an unembedded sentence may be merely entailed, merely presupposed, or both entailed and presupposed.

other words, the assumption that the CC of factive predicates is presupposed is widely entrenched. It is these analyses that motivated the current investigation: it is crucial to understand which predicates they apply to.

As illustrated, Kiparsky & Kiparsky 1970 characterized presupposed content as content that is not part of the main assertion and diagnosed it using projection. This view of presuppositions has been standard for a long time: for instance, in Beaver and Geurts's (2014) article in the *Stanford encyclopedia of philosophy*, presupposed content is characterized as content that is 'taken for granted, rather than being part of the main propositional content of a speech act' (abstract), and it notes that '[p]rojection from embeddings, especially negation, is standardly used as a diagnostic for presupposition' (§2). But the picture is more complex, because presuppositions are not the only backgrounded content that projects: conventional implicatures, like those contributed by appositives, are also backgrounded, projective content (see e.g. Chierchia & McConnell-Ginet 1990, Potts 2005). Thus, projection alone does not suffice to identify presupposed content.

There are, however, properties on which presuppositions and conventional implicatures differ, and we can use those properties to distinguish the two kinds of projective content. One such property is Potts's (2005) antibackgrounding requirement: conventional implicatures, unlike presuppositions, are unacceptable in a context in which they are part of the common ground. To illustrate, assume that the content of 6, that Lance Armstrong survived cancer, is part of the common ground. The continuation in 6a is unacceptable because this content is conventionally implicated by the appositive *a cancer survivor*. The continuation in 6b, by contrast, is acceptable: here, the same content is presupposed, as the CC of *know*.

- (6) Lance Armstrong survived cancer. (Potts 2005:34)
 - a. #When reporters interview Lance, a cancer survivor, he often talks about the disease.
 - b. And most riders know that Lance Armstrong is a cancer survivor.

Like the CC of *know*, illustrated in 6b, the CCs of clause-embedding predicates more generally are acceptable when their content is already part of the common ground: for instance, the example in 6b is acceptable when *know* is replaced with *say*, *believe*, *take into account*, or *confirm*. Thus, if a CC projects, we can assume that it is presupposed rather than conventionally implicated, because it does not exhibit the antibackgrounding requirement typical of conventional implicatures. Consequently, even though not every projective content is a presupposition, we can continue to assume that projective CCs are presupposed, as is standard practice in the field.⁴

There does not, to date, exist systematic empirical evidence that projection distinguishes factive predicates from others. To the contrary: there are hints in the literature that cast doubt on this assumption, which makes an investigation into the role of projection in classifying predicates a pressing issue. As a first hint, note that Kiparsky & Kiparsky 1970 identified not only factive and nonfactive predicates but also predicates that are neither factive nor nonfactive, including *acknowledge*, *remember*, and *announce*. Such predicates do not specify 'whether their complement sentences represent presuppositions by the speaker' (Kiparsky & Kiparsky 1970:164), and they 'can be

⁴ One exception is Anand and Hacquard (2014), who suggested (pp. 74f.) that even though the CC of predicates like *acknowledge*, *admit*, and *confirm* projects, the CC is not presupposed. They did not, however, offer a diagnostic for distinguishing projective CCs that are presupposed from projective CCs that are not. Ultimately, their argument against classifying such predicates as factive rests on the observation that their CC is not entailed.

used with or without presupposition of the truth of their complement sentence' (Karttunen 1971a:340). Thus, the CCs of these predicates, which we refer to as OPTIONAL FACTIVE, can be presupposed and, therefore, can project at least occasionally. It is an open, empirical question whether factive and optionally factive predicates are distinguishable by how regular the projection behavior of their CCs is.⁵

A second hint that casts doubt on the assumption that projection distinguishes factive predicates from others comes from Karttunen's (1971b:64) observation that there are purportedly factive predicates that 'lose their factivity' in polar questions and the antecedents of conditionals; he dubbed this subclass of predicates 'semifactive'. To illustrate, consider the negated examples in 7a, the polar questions in 7b, and the conditionals in 7c. Karttunen (1971b) suggested that 'all the examples in [7a] presuppose that John had not told the truth' (p. 63) and that a speaker who utters 7b or 7c with *regret* presupposes the CC: that is, the CC projects over negation, the polar question, and the antecedent of the conditional. However, he pointed out, a speaker who utters the variants of 7b and 7c with *realize* or *discover* does not necessarily presuppose the CC: the predicates 'permit both factive and nonfactive interpretations' (p. 63): that is, the CC may but need not project. Accordingly, *realize* and *discover*, but also *find out*, *see*, and *notice*, were taken to be semifactive, not properly factive, in contrast to predicates like *regret*, *forget*, *resent*, and *make clear*.

(7) Karttunen 1971b:63f.

- a. John didn't {regret/realize/discover} that he had not told the truth.
- b. Did you {regret/realize/discover} that you had not told the truth?
- c. If I {regret/realize/discover} later that I have not told the truth, I will confess it to everyone.

The observation that some purportedly factive predicates are merely semifactive again gives rise to the question of whether factive predicates are distinguishable from other predicates by how regular the projection behavior of their CCs is.⁶

As a third hint, note that not even the CCs of predicates assumed to be truly factive are necessarily presupposed. Consider *be aware* and *know*, which Karttunen (1971b) took to be truly factive: their CCs do not project in the naturally occurring examples in 8 from Beaver 2010. The observation that the CCs of factive predicates need not project and that the CCs of optionally factive predicates may project further increases the need to investigate whether projection distinguishes factive predicates from others.

- (8) a. Mrs. London is not aware that there have ever been signs erected to stop use of the route, nor that there ever has been any obstruction to stop use of the route. (Beaver 2010:83)
- b. Perhaps she knows that she really is to be married, and she, too, is now sad at the end of childhood. (Beaver 2010:86)
- c. If Susan knows that her eyes are dark brown, then A) She believes her eyes are dark brown B) Her belief that her eyes are dark brown is justified C) This belief is true D) All of the above E) None of the above. (Beaver 2010:84)

⁵ According to Kiparsky & Kiparsky 1970 and Karttunen 1971a, optionally factive predicates include *anticipate*, *acknowledge*, *suspect*, *report*, *emphasize*, *announce*, *admit*, *deduce*, and *remember*; this last predicate is taken to be factive in other research (e.g. Simons 2007, Kastner 2015, Abrusán 2016, Karttunen 2016, Aravind & Hackl 2017, Cremers 2018); in Karttunen & Zaenen 2005 and Karttunen 2016, *acknowledge*, *admit*, and *confess* are considered factive; *announce* is considered nonfactive in Karttunen & Zaenen 2005.

⁶ Contrary to what Karttunen 1971b assumed, nonprojection of the CC of semifactive predicates is possible not just in polar questions and conditionals (Beaver 2010).

Finally, some experimental results suggest that projection does not distinguish factive predicates from others. Tonhauser et al. 2018 investigated the projection of the CC of twelve clause-embedding predicates in polar questions, including ten canonically factive and two optionally factive predicates, namely *establish* and *confess*.⁷ As expected, Tonhauser et al. found that the CC of the twelve predicates investigated was projective and that the CC of the ten canonically factive predicates was more projective than that of the two optionally factive ones. However, they also observed significant projection differences among the canonically factive predicates, which gives rise to the question of whether projection would continue to distinguish factive from optionally factive predicates if a greater variety of optionally factive predicates were investigated. The findings of de Marneffe et al. (2019), who investigated the projection of the CC of forty-five clause-embedding predicates in naturally occurring data, suggest that this is not the case: although the CCs of canonically factive predicates were among the most projective, they were not categorically more projective than those of other predicates. We present these findings in detail in §2.3. Hints like these motivate a systematic investigation of whether projection identifies a coherent class of factive predicates.

ENTAILED. On the definition in 3b, the CC of factive predicates is presupposed and entailed: that is, for a predicate to be factive, its CC must project and also be an indefeasible inference from an affirmative matrix sentence with the predicate. The literature reveals disagreements for many predicates about whether the CC is entailed. Schlenker (2010:139), for instance, took the CC of *inform* to be entailed and *inform* to be a factive predicate: he assumed that Mary being pregnant necessarily follows from 9a. Anand and Hacquard (2014:76), however, pointed out that *inform* can be modified by *falsely*, as in the naturally occurring example in 9b, which they took to show that the CC does not necessarily follow from a true affirmative matrix sentence with *inform*; according to them, *inform* is not factive. Similarly, Swanson (2012) took the CC of *establish* to be entailed, but the naturally occurring example in 10 again suggests that the CC does not necessarily follow from affirmative matrix sentences with *establish*.

- (9) a. Mary has informed her parents that she's pregnant.
b. Doctors falsely informed my parents that I was dying.⁸
- (10) Diplomacy was repudiated, but with the media's help it was falsely established that the 'villain' was refusing to negotiate.⁹

Emotive predicates like *regret*, *be glad*, or *be annoyed* are also often taken to be factive, following Kiparsky & Kiparsky 1970 and Karttunen 1971b. This assumption has been challenged based on data like 11, where the truth of the CC does not follow from the affirmative matrix sentences.

- (11) a. Mary, who was under the illusion that it was Sunday, was glad that she could stay in bed. (Klein 1975, as cited in Gazdar 1979a:122 and Heim 1992, n. 37)
- b. Falsely believing that he had inflicted a fatal wound, Oedipus regretted killing the stranger on the road to Thebes. (Klein 1975)
- c. John wrongly believes that Mary got married, and he is enraged that she is no longer single. (adapted from Egré 2008, based on Schlenker 2003a)

⁷ We assume that these predicates may be taken to be optionally factive given work that does not find speakers to presuppose the truth of the CC of these predicates (e.g. Wyse 2010, Swanson 2012, Karttunen 2016).

⁸ <https://www.quora.com/Has-a-doctor-ever-falsely-informed-you-that-you-re-dying>

⁹ <https://books.google.de/books?id=BnU1qPRO-34C>, p. 66

The existence of examples like those in 11 has led some authors, including Klein (1975), Giannakidou (1998), Schlenker (2003b), and Egré (2008), to conclude that emotive predicates are not factive: a speaker who uses one of these predicates merely presupposes that the subject of the attitude believes the truth of the CC (Heim 1992). Consistent with this position is a proposal by Karttunen (2016) according to which the CC is not entailed but is taken to follow from utterances of affirmative matrix sentences as long as there is no indication that the beliefs of the speaker/writer and the subject of the attitude are not aligned; along these lines see also Djärv 2019. Other works have maintained that emotive predicates are factive despite examples like 11 (e.g. Gazdar 1979a, Abrusán 2011, Anand & Hacquard 2014): they hypothesize that examples like those in 11 are licensed by a kind of free indirect discourse—that is, these examples do not show that the CC of such predicates is not entailed (or presupposed).

There is similar disagreement about the cognitive predicate *know*: the examples in 12 show that the truth of the CC does not necessarily follow from true affirmative matrix sentences with *know*. Abrusán (2011:515) assumed that these examples do not show that *know* is not factive because ‘they are understood as reporting a firm belief or feeling of “knowledge” on somebody’s part’, maintaining that these examples ‘are analogous to the free indirect discourse cases’ discussed above.

- (12) a. Everyone knew that stress caused ulcers, before two Australian doctors in the early 80s proved that ulcers are actually caused by bacterial infection.
(Hazlett 2010:501)
- b. John suffers from paranoia. He falsely believes that the police is spying on him and what is more he knows they are listening to his phone calls.
(Abrusán 2011:514)
- c. Most educated people already have opinions about language. They know that it is man’s most important cultural invention ... I will try to convince you that every one of these common opinions is wrong! (Pinker 1994:17f.)¹⁰

In short, there is disagreement for many predicates about whether their CC is entailed. Under the definition in 3b, the uncertainty about whether the CC is entailed means that it is unclear whether these predicates are factive. There has not, to date, been a systematic investigation of which clause-embedding predicates have entailed CCs.¹¹

1.3. THE GOAL OF THIS ARTICLE AND THE INVESTIGATED CLAUSE-EMBEDDING PREDICATES. As illustrated, there is no consensus in the literature about which predicates are factive. This lack of consensus is due to differences in how factive predicates are defined, and it is compounded by uncertainty about whether projection categorically distinguishes factive predicates from others and by disagreement about whether the CC of particular predicates is entailed. This article investigates the projection and entailment of the CCs of twenty clause-embedding predicates, with the goal of understanding whether these two properties identify a coherent class of factive predicates. The predicates we investigate are listed in 13 with the categories they are typically taken to fall into. Specifically, the five canonically factive predicates in 13a include the cognitive

¹⁰ We thank Manfred Krifka (p.c.) for pointing out this example to us.

¹¹ White and Rawlins’s (2018) MegaVeridicality data set and Ross and Pavlick’s (2019) VerbVeridicality data set include entailment and projection ratings, but the authors did not present by-predicate analyses. We compare the MegaVeridicality and VerbVeridicality ratings to our data in §2.3 and §3.5. Djärv et al. (2016) set out to investigate whether emotive and cognitive predicates differ in whether the CC is entailed. However, because the predicates in their stimuli were realized in polar questions (e.g. *Is Maria aware that Mike is moving back to Chicago?*), it is not clear that the observed differences bear on whether the CC is entailed.

predicate *know*, the change-of-state predicates *discover* and *reveal*, the sensory predicate *see*, and the emotive predicate *be annoyed*. The six nonfactive predicates in 13b include four NONVERIDICAL NONFACTIVE predicates *pretend*, *suggest*, *say*, and *think*, whose CC is typically taken to be neither presupposed nor entailed, as well as the two VERIDICAL NONFACTIVE predicates *be right* and *demonstrate*, whose CC is typically taken to be entailed but not presupposed.¹² The nine optionally factive predicates in 13c include *acknowledge*, *admit*, and *announce* (see Kiparsky & Kiparsky 1970) as well as *confirm*, *inform*, *confess*, *establish*, *hear*, and *prove*.¹³

- (13) Twenty clause-embedding predicates
 - a. canonically factive: *be annoyed*, *discover*, *know*, *reveal*, *see*
 - b. nonfactive:
 - i. nonveridical nonfactive: *pretend*, *say*, *suggest*, *think*
 - ii. veridical nonfactive: *be right*, *demonstrate*
 - c. optionally factive: *acknowledge*, *admit*, *announce*, *confess*, *confirm*, *establish*, *hear*, *inform*, *prove*

Under the definition in 3a, according to which the CC of factive predicates is presupposed, we expect to see a categorical difference in projection between canonically factive predicates on the one hand, and optionally factive and nonfactive predicates on the other. Section 2 reports two experiments designed to investigate the projection of the CCs of these twenty clause-embedding predicates, the results of which allow us to assess whether projection identifies a class of factive predicates. Under the definition in 3b, the CC of factive predicates is both presupposed and entailed. Section 3 reports four experiments designed to investigate whether the CC of the twenty clause-embedding predicates is entailed; we expect the CC of canonically factive and veridical nonfactive predicates to be entailed. These results allow us to assess whether projection and entailment jointly identify a class of factive predicates.

2. EXPERIMENT 1: PRESUPPOSITION. Experiments 1a and 1b were designed to investigate which of the CCs of the twenty clause-embedding predicates in 13 are presupposed.¹⁴

The experiments employed the CERTAIN THAT diagnostic for projection (see also e.g. Tonhauser 2016, Djärv & Bacovcin 2017, Stevens et al. 2017, Lorson 2018, Tonhauser et al. 2018, de Marneffe et al. 2019, Mahler 2019, 2020): on this diagnostic, participants are presented with an utterance in which the content to be investigated occurs in an en-

¹² Anand & Hacquard 2014 assumed *demonstrate* to be veridical, in contrast to Anand et al. 2019. This latter work also took *reveal* to be nonfactive, in contrast to, for instance, Egré 2008, Wyse 2010, and Tonhauser et al. 2018.

¹³ Different categories have been assumed for these six predicates. Anand and Hacquard (2014) took *confirm* and *inform* to not be factive, while Schlenker (2010) took *inform* to be factive. For *prove*, White and Rawlins (2018) suggested that it is veridical and nonfactive, but Anand and Hacquard (2014) took it to be nonveridical and nonfactive. Swanson (2012) took *confess* to be factive, while Karttunen (2016) only took it to commit the speaker to the subject of the attitude being committed to the CC, and Wyse (2010) listed it under the nonfactive predicates. Swanson (2012) took *establish* to be nonfactive, but Wyse (2010) listed it under the factive predicates. We also included *hear*: even though it is often considered a factive predicate (e.g. Beaver 2010, Anand & Hacquard 2014), *hear* also has a nonfactive reportative evidential sense (see e.g. Anderson 1986, Simons 2007), especially when combined with complements that describe events that cannot be auditorily perceived, as in our experiments.

¹⁴ The experiments, data, and R code for generating the figures and analyses of the experiments reported in this article are available at <https://github.com/judith-tonhauser/projective-probability>. All experiments were conducted with IRB approval. Experiments 1b, 2b, and 3b were preregistered: <https://osf.io/cxq47>.

tailment-canceling environment, like a polar question, as illustrated in 14 for the content of the appositive, that Martha's new car is a BMW.¹⁵

- (14) Patrick asks: *Was Martha's new car, a BMW, expensive?*

To assess projection, participants are asked whether the speaker is certain of the content; for instance, in 14, participants are asked whether Patrick is certain that Martha's new car is a BMW. If a participant takes the speaker to be certain of the content, we assume that the content projects, that is, it is presupposed by the speaker; if a participant does not take the speaker to be certain of the content, we assume that the content does not project, that is, it is not presupposed.

Following Tonhauser et al. 2018, experiment 1a implemented the ‘certain that’ diagnostic with a gradient response scale: participants gave their responses on a slider marked ‘no’ at one end and ‘yes’ at the other. We assume that the closer to ‘yes’ a participant’s response is, the more the speaker is certain of the content, that is, the more projective the content is. In other words, we assume that gradient certainty ratings reflect gradience in the speaker’s commitment to the truth of the content.¹⁶ As discussed in Tonhauser et al. 2018, a second interpretation of gradient certainty ratings is that they reflect a participant’s degree of belief in the speaker’s (binary) commitment to the truth of the content. While we adopt the first interpretation in our discussion, we remain agnostic about the interpretation of gradient certainty ratings: either interpretation gives rise to the expectation, on the definition of factive predicates in 3a, that certainty ratings should be categorically higher for factive predicates than for optionally factive and non-factive ones.

Experiment 1b was identical to experiment 1a but employed a two-alternative forced-choice task where participants categorically responded ‘yes’ or ‘no’.

2.1. EXPERIMENT 1A: GRADIENT PROJECTION RATINGS.

METHODS.

Participants. Three hundred participants with US IP addresses and at least 99% of previously approved human intelligence tasks (HITs) were recruited on Amazon’s Mechanical Turk platform. They were paid \$0.75.

Materials. Experiment 1a investigated the projection of the CC of the twenty clause-embedding predicates in 13. The clausal complements of the twenty predicates were realized by twenty clauses (provided in §A of the online supplementary materials), for a total of 400 predicate + clause combinations.¹⁷ These 400 predicate + clause combinations were realized as polar questions by combining them with random proper-name subjects, as shown in the sample target stimuli in 15, which realize the complement clause *Julian dances salsa*. Eventive predicates, like *discover* and *hear*, were realized in the past tense, and stative predicates, like *know* and *be annoyed*, were realized in the present tense. The indirect object of *inform* was realized by the proper name *Sam*. The speaker of the target stimuli was realized by a proper name and displayed in boldface.

¹⁵ For other diagnostics for projection see, for example, Smith & Hall 2011, Xue & Onea 2011, and Tonhauser et al. 2013; see also the discussion in Tonhauser et al. 2018.

¹⁶ Strictly speaking, gradient certainty ratings reflect gradience in the degree to which participants PERCEIVE the speaker to be committed to the truth of the content. That is, as in any experiment, the quantity of interest, in this case speaker commitment, is only indirectly measured.

¹⁷ The supplementary materials referenced here and elsewhere are available at <http://muse.jhu.edu/resolve/160>.

The proper names that realized the speakers, the subjects of the clause-embedding predicates, and the subjects of the complement clauses were all unique.

- (15) a. **Carol asks:** Did Sandra discover that Julian dances salsa?
- b. **Paul asks:** Does Edward think that Julian dances salsa?

To assess whether participants were attending to the task, the experiment also included six control stimuli, whose main clause content was investigated. For instance, in the sample control stimulus in 16, participants were asked whether the speaker was certain that Zack is coming to the meeting tomorrow. We expected participants to rate the projection of these contents as very low because speakers are typically not committed to content they are asking about. For the full set of control stimuli see supplement §B.

- (16) Is Zack coming to the meeting tomorrow?

Each participant saw a random set of twenty-six stimuli: each set contained one target stimulus for each of the twenty clause-embedding predicates (each with a unique complement clause) and the same six control stimuli. Trial order was randomized.

Procedure. Participants were told to imagine that they are at a party and that, on walking into the kitchen, they overhear somebody ask a question. Participants were asked to rate whether the speaker was certain of the CC. They gave their responses on a slider marked ‘no’ at one end (coded as 0) and ‘yes’ at the other (coded as 1), as shown in Figure 1.

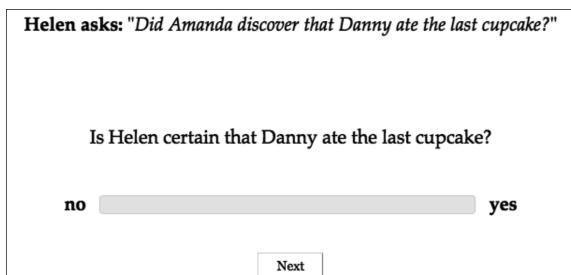


FIGURE 1. A sample trial in experiment 1a.

Data exclusion. The data from thirty-four participants were excluded, based on self-declared nonnative speaker status and other criteria given in supplement §C. The data from 266 participants (ages: 20–71, median: 36; 118 female, 143 male, 2 other, 3 undeclared) were analyzed.

RESULTS AND DISCUSSION. Figure 2 shows the mean certainty ratings for the target stimuli by predicate, as well as for the main clause stimuli (abbreviated ‘MC’), in increasing order from left to right.¹⁸ The mean certainty ratings were largely consistent with impressionistic judgments reported in the literature. First, the ratings for main clause content were lowest overall, as expected for nonprojective content. Second, the ratings for canonically factive predicates were among the highest overall, suggesting comparatively high projection of the CCs. Third, the mean certainty ratings of many optionally factive predicates were lower than those of many factive predicates and higher

¹⁸ The width of the violin at any given point represents the proportion of individual participants’ ratings located at that point. This figure, like others in the article, is presented in color in the electronic versions of this article, but in grayscale in the print version. Color versions of the figures are also available open access along with the supplementary materials at <http://muse.jhu.edu/resolve/160>.

than those of main clauses as well as of nonveridical nonfactives. However, Fig. 2 also shows that the CCs of the five canonically factive predicates were not categorically more projective than the CCs of the optionally factive predicates, contrary to what is expected under the first definition of factive predicates. Specifically, the CCs of the optionally factive *acknowledge*, *hear*, and *inform* were at least as projective as the CCs of *reveal* or *discover*. This result suggests that projection does not categorically distinguish factive predicates from optionally factive and nonfactive ones.

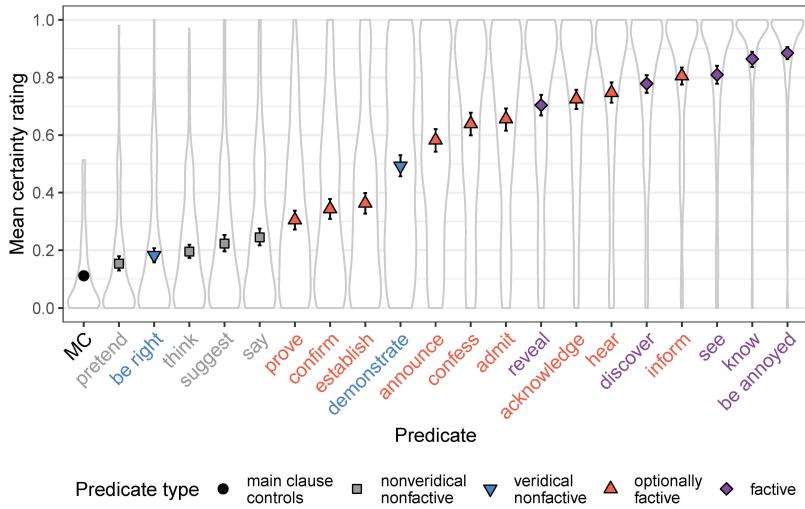


FIGURE 2. Mean certainty ratings by predicate in experiment 1a. Error bars indicate 95% bootstrapped confidence intervals. Violin plots indicate the kernel probability density of the individual participants' ratings.

In view of this result, one might ask whether the categorization of clause-embedding predicates assumed in 13 is correct: perhaps there is a different place where a line can be drawn between factive predicates and others based on the projection of the CCs? The mean certainty ratings visualized in Fig. 2 suggest that this cannot be done in a nonarbitrary way. For instance, one might consider all predicates factive whose CC is at least as projective as that of *reveal*. This would mean that canonically factive predicates are factive, as well as *acknowledge*, *hear*, and *inform*. There is, however, no principled reason why the projection of the CC of *reveal* should be the cutoff for factivity. A different approach would be to pick an arbitrary mean certainty rating, say 0.8, and to categorize predicates whose CC has a mean certainty rating of at least 0.8 as factive and those whose CC has a mean certainty rating below 0.8 as optionally factive. However, this is as unprincipled an approach as the previous one—one could just as well pick a threshold of 0.85 or 0.75, with the result that different predicates would be considered factive. For instance, *discover* (mean: 0.78) would not count as factive with a threshold of 0.8, but would with a threshold of 0.75, and *see* (mean: 0.81) would count as factive with a threshold of 0.8, but not with a threshold of 0.85.¹⁹

¹⁹ Choosing an arbitrary numeric threshold for factivity is also complicated by the fact that there is by-experiment variation in mean certainty ratings. For instance, the mean certainty ratings in experiment 1a were lower, overall, than those in Tonhauser et al.'s (2018) experiments 1a,b. For example, the CC of *be annoyed*, which was the most projective in both sets of experiments, received mean certainty ratings of 0.96 and 0.92 in

The possibility of categorizing clause-embedding predicates by the projection of the CC is further called into question by the fact that the CC of all twenty predicates investigated is projective, albeit to varying degrees, compared to the nonprojective main clause controls. This was established by fitting a Bayesian mixed-effects beta regression model with weakly informative priors using the ‘*brms*’ package (Bürkner 2017) in R (R Core Team 2016). The model predicted the certainty ratings from a fixed effect of predicate (with treatment coding and ‘main clause control’ as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. This random-effects structure captures random variability in projection between participants and between items, where an item is a unique combination of a predicate and a complement clause. A beta regression model estimates the mean of the outcome distribution (like a linear regression model).²⁰ We thus obtain, for each predicate, a 95% credible interval for the mean rating for that predicate: this allows us to identify whether the mean ratings of the predicate and of the main clause controls differ. Supplement §D motivates the use of beta regression over linear regression, provides a brief primer on how to interpret Bayesian mixed-effects beta regression models, and reports the full model output.

According to the beta regression model, the estimated mean for each predicate was higher than that of the main clause controls: that is, the 95% credible intervals for the estimated adjustment to the main clause control mean did not contain 0 for any predicate. This result suggests that the CC of each of the twenty predicates is projective compared to nonprojective main clause content, albeit to different degrees.²¹ Thus, to distinguish factive predicates from others in this set of twenty predicates, one would need to distinguish one group of projective CCs from another group of projective CCs.

In sum, the results of experiment 1a suggest that projection does not identify a coherent class of factive predicates. It is possible, however, that asking participants to assess projection with a gradient response scale encouraged them to provide noncategorical answers and that the observed projection variability is an artifact of the task. If so, categorically distinguishing factive predicates from others may be possible if participants were asked to assess projection with a categorical response task. To assess this possibility, we conducted experiment 1b, which was identical to experiment 1a but employed a two-alternative forced-choice task.

Tonhauser et al.’s experiments 1a and 1b, respectively, but only a 0.86 in our experiment 1a. This difference may be due to the fact that Tonhauser et al. primarily investigated highly projective content, whereas our experiment 1a included a wide range of less projective content: it is possible that participants’ certainty ratings were influenced by the overall projection of the contents investigated.

²⁰ Beta regression models also estimate a second parameter, namely the precision, which is a measure of dispersion: the greater the precision, the more concentrated the ratings are around the mean. In this article, we rely on the estimated means to identify whether the CCs of the clause-embedding predicates differ from the relevant controls. Both the estimated mean and precision for each predicate are reported in the full model output tables in supplement §D.

²¹ A Bayesian mixed-effects linear regression with the same fixed- and random-effects structure yielded qualitatively identical results, except that the contrast between *pretend* and the main clause controls was only marginally significant. See the Github repository mentioned in n. 14 for the model code.

2.2. EXPERIMENT 1B: CATEGORICAL PROJECTION RATINGS.

METHODS.

Participants. Six hundred participants with US IP addresses and at least 99% of previous HITs approved were recruited on Amazon’s Mechanical Turk platform. They were paid \$1.²²

Materials and procedure. The materials and procedure were identical to those of experiment 1a, with the exception of the task: participants responded ‘yes’ (coded as 1) or ‘no’ (coded as 0) to the question of whether the speaker is certain of the relevant content, as shown in Figure 3.

Matthew asks: "Did Anna prove that Charley speaks Spanish?"

Is Matthew certain that Charley speaks Spanish?

No Yes

Next

FIGURE 3. A sample trial in experiment 1b.

Data exclusion. The data from 164 participants were excluded according to the criteria given in supplement §C. The data from 436 participants (ages: 18–81, median: 37; 207 female, 223 male, 2 other, 4 undeclared) were analyzed.

RESULTS AND DISCUSSION. Figure 4 shows the proportion of ‘yes’ ratings (indicating projection) for the predicates and the main clause stimuli, in increasing order from left to right. Jittered gray dots indicate the individual participants’ ratings.²³ The order of predicates according to the projection of their CC in experiment 1b is strikingly similar to that of experiment 1a, as evidenced by the very high Spearman rank correlation of 0.983 (see supplement §E for a visualization and details on this correlation).

The critical results of experiment 1a were replicated. First, the CCs of canonically factive predicates were not categorically more projective than the CCs of the optionally factive predicates: the CCs of *acknowledge*, *hear*, and *inform* were again at least as projective as those of *reveal* or *discover*. Second, a nonarbitrary line between canonically factive and optionally factive predicates cannot be drawn: the CCs of all twenty predicates were projective compared to the nonprojective main clause controls.²⁴ This was established by fitting a Bayesian mixed-effects logistic regression model with weakly informative priors using the brms package in R. The model predicted the log-odds of a ‘yes’ over a ‘no’ response from a fixed effect of predicate (with treatment coding and ‘main clause’ as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. The estimated log-odds of a ‘yes’ response was greater for each predicate than that of the main

²² Experiments 1a, 2a, and 3a were run in 2017 and 2018, whereas experiments 1b, 2b, and 3b were run in 2019. Participants in the later experiments were paid more to reflect an increased minimum wage in California.

²³ Recall that participants gave binary certainty ratings, with ‘yes’ coded as 1 and ‘no’ coded as 0. To increase legibility, these ratings are jittered vertically and horizontally in Fig. 4 and other figures like it. Consequently, the individual certainty ratings should not be read against the y-axis.

²⁴ According to Spector & Egré 2015:1739, it is not clear whether the CC of *say* is projective. The results of our experiments 1a,b suggest that it is weakly projective.

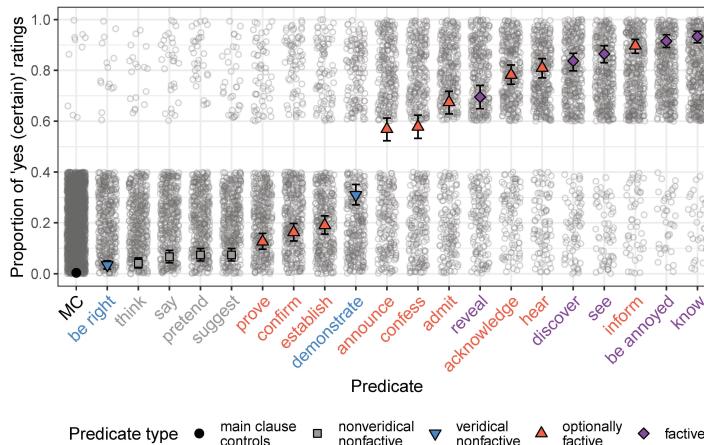


FIGURE 4. Proportion of ‘yes’ ratings by predicate in experiment 1b. Error bars indicate 95% bootstrapped confidence intervals. Vertically and horizontally jittered light gray dots indicate individual participants’ responses of ‘yes’ (coded as 1) and ‘no’ (coded as 0).

clause controls: that is, the 95% credible intervals for the estimated coefficients did not contain 0 for any predicate. See supplement §D for model details. Thus, like experiment 1a, the results of experiment 1b do not provide support for the definition of factive predicates in 3a, according to which factive predicates are categorically distinguished from optionally factive and nonfactive predicates by the projection of the CC.

2.3. INTERIM SUMMARY AND CONVERGING EVIDENCE. The results of experiments 1a,b suggest that projection does not categorically distinguish canonically factive predicates from others: contrary to what is expected on definition 3a, we observed that the CCs of predicates standardly classified as factive are not categorically more projective than the CCs of all other predicates. We now discuss converging evidence to this effect from three additional data sets: the CommitmentBank (de Marneffe et al. 2019), the MegaVeridicality data set (White & Rawlins 2018, White et al. 2018), and the VerbVeridicality data set (Ross & Pavlick 2019).²⁵

The CommitmentBank is a collection of 1,200 naturally occurring English discourses from the *Wall Street Journal*, the British National Corpus, and the Switchboard corpus: each discourse consists of a target sentence with a clause-embedding predicate embedded in an entailment-canceling operator and up to two preceding context sentences, as illustrated in 17, with *know* in the target sentence embedded under negation.

- (17) What fun to hear Artemis laugh. She’s such a serious child. I didn’t know she had a sense of humor. (de Marneffe et al. 2019:109)

Annotators recruited on Amazon’s Mechanical Turk platform provided certainty ratings for the CCs of the target sentences on a seven-point Likert scale labeled at three points: speaker/author is certain that the CC is true (coded as 3), speaker/author is not certain whether the CC is true or false (coded as 0), and speaker/author is certain that the CC is false (coded as -3); the lower half of the scale was included to differentiate speaker/

²⁵ We replotted the data, obtained at <https://github.com/mcdm/CommitmentBank>, <http://megaattitude.io>, and https://github.com/alexisjihyross/verb_veridicality, respectively. For the analysis files, see the GitHub repository mentioned in n. 14.

author noncommitment to the CC from speaker/author commitment to the falsity of the CC (as may arise in, for instance, *I don't know that I agree*). De Marneffe et al. (2019) assumed, for certainty ratings above 0, that the higher the certainty rating, the more projective the CC. Figure 5 shows the mean certainty ratings for the CCs in the 982 discourses with forty-five nonveridical nonfactive, optionally factive, and factive clause-embedding predicates (target sentences embedded under nonepistemic modals were excluded; see de Marneffe et al. 2019:§3). In line with the results of experiments 1a and 1b, projection of the CC does not categorically distinguish canonically factive predicates from optionally factive and nonfactive ones: for instance, the CCs of the optionally factive *foresee*, *confess*, *show*, *tell*, and *accept* are at least as projective as those of some canonically factive predicates.²⁶

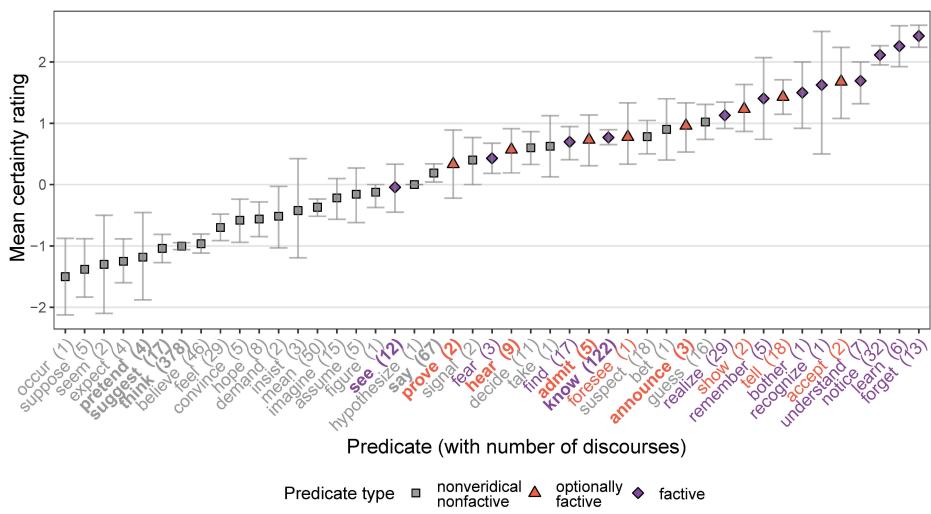


FIGURE 5. Mean certainty ratings by predicate, with number of discourses in parentheses, for 982 discourses in the CommitmentBank (de Marneffe et al. 2019). Error bars indicate 95% bootstrapped confidence intervals. Predicates included in our experiments 1a,b shown in bold.

A second piece of converging evidence that projection does not identify a class of factive predicates comes from the VerbVeridicality data set. This data set includes 859 naturally occurring indicative matrix sentences with a total of seventy-eight clause-embedding predicates and their clausal complements from the MultiNLI (Williams et al. 2018) corpus. To diagnose projection of the CC, negated variants of the 859 matrix sentences were created; an example is given in 18. For each negated variant, three participants rated on a five-point Likert scale whether the CC is true given the target sentence; one endpoint of the scale was labeled ‘definitely true’ (coded as 2) and the other ‘definitely not true’ (coded as -2).

- (18) The GAO has not indicated that it is unwilling to compromise.

(Ross & Pavlick 2019:2234)

Figure 6 plots the mean projection ratings for the CCs of the seventy-eight predicates in the VerbVeridicality data set, with labels for the fifteen predicates from our experiments 1a,b that are included (*see* and *saw*, which are both included in the VerbVeridical-

²⁶ De Marneffe et al. (2019) found that a model that predicts certainty ratings from factivity captures less of the variance than a model that predicts certainty ratings from predicate lemma. This result is replicated in our experiments 1a,b, as discussed in §4.

ity data set, are plotted separately). As shown, projection does not categorically distinguish canonically factive from optionally factive and nonfactive predicates: the CCs of all of our optionally factive predicates are at least as projective as those of some canonically factive predicates.

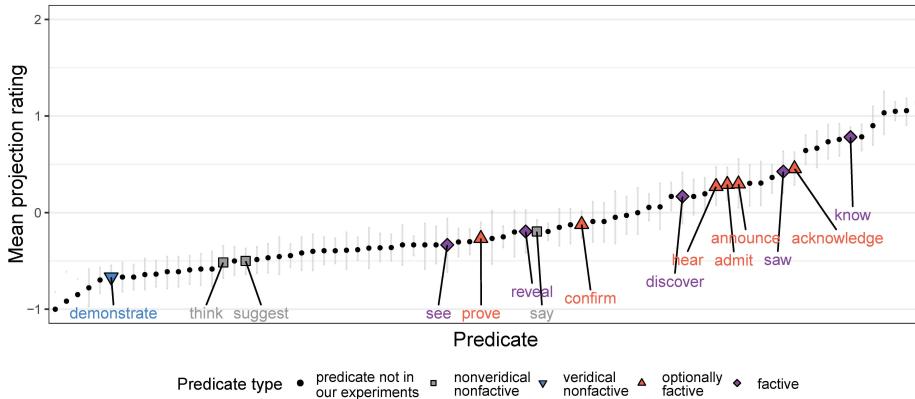


FIGURE 6. Mean projection ratings by predicate, with 95% bootstrapped confidence intervals, for the seventy-eight predicates in the VerbVeridicality data set (Ross & Pavlick 2019), with labels for the fifteen predicates featured in our experiments.

A third piece of converging evidence comes from the MegaVeridicality data set, which contains projection ratings for the CCs of 517 English clause-embedding predicates. The stimuli that participants rated consisted of combinations of these predicates with what White and Rawlins (2018) referred to as ‘low content arguments’, as shown in 19 for *know*. The predicates were embedded under negation in stimuli like 19a, and under negation, in the antecedent of a conditional, and in a question in stimuli like 19b. To assess projection, participants were asked to respond to the question *Did that thing happen?* for stimuli like 19a and to respond to the question posed by stimuli like 19b. The response options were ‘yes’, ‘maybe or maybe not’, and ‘no’.

- (19) a. Somebody didn’t know that a particular thing happened.
 b. If somebody didn’t know that a particular thing happened, did that thing happen?

The CCs of each of the predicates in the MegaVeridicality data set received between twenty-nine and sixty projection ratings (mean: thirty-two). To plot the ratings, we coded a ‘yes’ response as 1 (the speaker is certain that the thing happened), a ‘maybe or maybe not’ response as 0 (the speaker is not certain whether the thing happened), and a ‘no’ response as -1 (the speaker is certain that the thing did not happen). Figure 7 displays the mean projection ratings for the 517 predicates in the MegaVeridicality data set, with labels for the nineteen predicates from our experiments that are included. As shown, projection of the CC does not categorically distinguish canonically factive predicates from others: for instance, the CCs of the optionally factive predicates *confess*, *acknowledge*, *admit*, *announce*, *hear*, and *inform* are at least as projective as those of some canonically factive predicates.

In sum, contrary to what is expected under the definition of factive predicates in 3a, the results of our experiments 1a and 1b as well as our meta-analyses of the CommitmentBank, the VerbVeridicality data set, and the MegaVeridicality data set suggest that projection of the CC does not identify a coherent class of factive predicates. The chal-

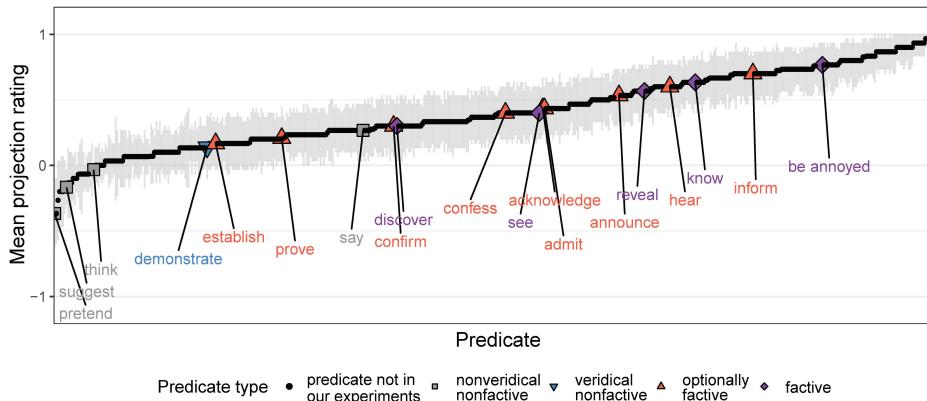


FIGURE 7. Mean projection ratings by predicate, with 95% bootstrapped confidence intervals, for the 517 predicates in the MegaVeridicality data set (White & Rawlins 2018, White et al. 2018), with labels for the nineteen predicates featured in our experiments.

length to definition 3a is compounded by the diversity of the empirical evidence: it comes from constructed and naturally occurring examples, from examples presented with and without a context, from examples with diverse and with minimal lexical content, from clause-embedding predicates embedded under different entailment-canceling operators, and from projection ratings collected through different response tasks.

The experiments described in the next section were designed to explore whether the definition of factive predicates in 3b, according to which presupposition and entailment jointly identify a class of factive predicates, is empirically supported. To this end, we collected entailment judgments for CCs of the same twenty predicates studied in experiments 1a,b.

3. EXPERIMENTS 2 AND 3: ENTAILMENT. Experiments 2 and 3 explored which of the CCs of the twenty clause-embedding predicates are entailed, to identify whether the definition of factive predicates in 3b identifies a coherent class. We assume the standard definition of entailment, according to which entailment is a binary, categorical relation between two contents: given sentences ϕ and ψ , the content of ϕ entails the content of ψ if and only if every world in which ϕ is true is also a world in which ψ is true. If there is at least one world in which ϕ is true but ψ is false, then ϕ does not entail ψ .

To establish whether the content of an indicative matrix sentence with a clause-embedding predicate entails the CC, we used two standard diagnostics for entailment, given in 20. According to the INFERENCE DIAGNOSTIC in 20a, the content of an indicative matrix sentence with a clause-embedding predicate entails the CC if and only if the CC definitely follows from a true statement of the matrix sentence. According to the CONTRADICTION DIAGNOSTIC in 20b, an entailment relationship holds if and only if an utterance of the indicative matrix sentence that is followed by a denial of the CC is contradictory.

- (20) Two diagnostics for entailment (see e.g. Chierchia & McConnell-Ginet 1990:§3.1)
 - a. Inference diagnostic (experiments 2a and 2b): ϕ entails ψ if and only if, if ϕ is true, then the truth of ψ definitely follows.
 - b. Contradictoriness diagnostic (experiments 3a and 3b): ϕ entails ψ if and only if a sentence of the form ϕ and not ψ is contradictory.

These two diagnostics for entailment were applied to the CCs of the twenty clause-embedding predicates in experiments 2a/2b and experiments 3a/3b, respectively: the two versions of each experiment differed in whether participants provided a gradient response (a) or a categorical one (b). Both sets of experiments included control stimuli for which the relevant contents stand in an entailment relation, that is, for which the relevant inference definitely follows (experiment 2) and which are definitely contradictory (experiment 3). We assume that the CC of a clause-embedding predicate is not entailed if responses to the CC are significantly different from responses to these control stimuli, and that the CC of a clause-embedding predicate is entailed if responses to the CC are not significantly different from responses to these control stimuli. Thus, identifying entailed CCs requires arguing from a null result. Given the standard classification of the twenty clause-embedding predicates in 13, we expect the CCs of canonically factive and veridical nonfactive predicates to be entailed, that is, to not be judged differently from the control stimuli.

3.1. EXPERIMENT 2A: GRADIENT RATINGS USING THE INFERENCE DIAGNOSTIC FOR ENTAILMENT. Experiment 2a investigated which of the CCs of the twenty clause-embedding predicates are entailed based on the inference diagnostic for entailment in 20a. Participants assessed on a gradient scale whether the CC follows from true statements of indicative matrix sentences with the twenty clause-embedding predicates.

METHODS.

Participants. Three hundred participants with US IP addresses and at least 99% of previous HITs approved were recruited on Amazon’s Mechanical Turk platform. They were paid \$0.75.

Materials. The target stimuli were 400 indicative matrix sentences. They were constructed by combining the twenty clause-embedding predicates with the twenty complement clauses (for 400 predicate + clause combinations, as in experiment 1) and a proper-name subject whose gender differed from the gender of the proper-name subject of the embedded clause. As illustrated in the sample stimuli in 21, these matrix sentences were presented to participants as true statements. For each of the 400 target stimuli, the inference that was tested was the inference to the CC: for instance, participants were asked about 21a whether it follows that Danny ate the last cupcake and about 21b whether it follows that Emma studied on Saturday morning.

- (21) a. **What is true:** Melissa knows that Danny ate the last cupcake.
- b. **What is true:** Jerry pretended that Emma studied on Saturday morning.

The experiment included eight control stimuli that were used to assess whether participants were attending to the task and had interpreted the task correctly, as well as to identify inference ratings for entailed content. There were four entailing control stimuli for which the tested inferences are entailments of the true statements: in the sample entailing control in 22, for instance, that Frederick solved the problem is an entailment of the true statement that Frederick managed to solve the problem. We expected the inference ratings for these four control stimuli to be at ceiling, and, as discussed above, we rely on the ratings for these stimuli to identify entailments. In contrast, the inferences tested with the four nonentailing control stimuli, like 23, definitely do not follow from these true statements: for instance, in 23, that Dana wears a wig does not follow from Dana watching a movie last night. We expected the inference ratings for these four control stimuli to be at floor. See supplement §B for the full set of control stimuli.

- (22) **What is true:** Frederick managed to solve the problem. (tested inference: Frederick solved the problem.)
- (23) **What is true:** Dana watched a movie last night. (tested inference: Dana wears a wig.)

Each participant saw a random set of twenty-eight stimuli: each set contained one target stimulus for each of the twenty predicates (each with a unique complement clause) and the same eight control stimuli. Trial order was randomized.

Procedure. Participants were told that they would read true statements and that they would be asked to assess whether a second statement follows from that statement. On each trial, participants read the true statement and were asked to respond to the question *Does it follow that ... ?*, with the complement clause realized by the complement clause of the clause-embedding predicate. Participants responded on a sliding scale from ‘definitely doesn’t follow’ (coded as 0) to ‘definitely follows’ (coded as 1), as shown in Figure 8.

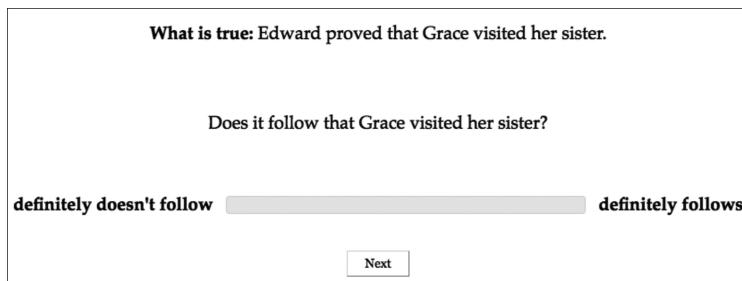


FIGURE 8. A sample trial in experiment 2a.

In order to familiarize participants with the task, they first responded to the two familiarization stimuli in 24, where the inference tested was the CC. Participants who rated 24a in the lower half of the sliding scale or 24b in the upper half of the sliding scale were given an explanation for why their answer was wrong. Participants could only advance to the twenty-eight stimuli if they gave a plausible rating to the two familiarization stimuli, that is, a rating in the upper half of the scale for 24a and in the lower half for 24b.

- (24) a. **What is true:** Drew is correct that Patty lives in Canada.
 b. **What is true:** Drew believes that Patty lives in Canada.

Data exclusion. The data from forty-one participants were excluded according to the criteria given in supplement §C. The data from 259 participants (ages: 19–69, median: 36; 132 female, 127 male) were analyzed.

RESULTS. Figure 9 shows the mean inference ratings for the target stimuli by predicate, in increasing order from left to right, as well as for the nonentailing and entailing controls. In line with impressionistic judgments reported in the literature, the mean inference ratings for most canonically factive predicates as well as the veridical nonfactive predicate *be right* were quite high, suggesting strong inferences to the truth of the CC. The mean inference ratings of the canonically factive predicate *reveal* and of the canonically veridical nonfactive predicate *demonstrate* were lower, suggesting comparatively weaker inferences to the truth of the CC.

To assess which predicates have entailed CCs, we fitted a Bayesian mixed-effects beta regression model with weakly informative priors using the brms package in R,

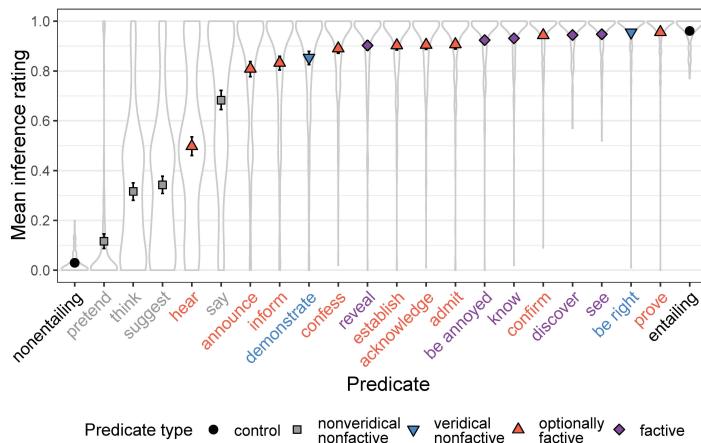


FIGURE 9. Mean inference ratings by predicate in experiment 2a, including the nonentailing and entailing controls. Error bars indicate bootstrapped 95% confidence intervals. Violin plots indicate the kernel probability density of the individual participants' ratings.

which predicted the inference ratings from a fixed effect of predicate (with treatment coding and 'entailing control' as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. According to this model, the estimated mean for each predicate except *prove* and *be right* was lower than that of the entailing controls: that is, the 95% credible intervals for the estimated adjustment to the intercept mean did not contain 0 for any predicate other than *prove* and *be right*. See supplement §D for the full model output.²⁷ Thus, the gradient inference ratings for the CCs of *be right* and *prove* are compatible with the assumption that their CCs are entailed. The inference ratings for the remaining predicates, including the canonically factive predicates *see*, *discover*, *know*, *be annoyed*, and *reveal* and the veridical nonfactive predicate *demonstrate*, were significantly lower, suggesting that these CCs are not entailed, contrary to assumption.

DISCUSSION. The results of experiment 2a have two unexpected implications. The first concerns the predicates for which experiment 2a found their CCs to be entailed, namely *be right* and *prove*. Given definition 3b, according to which a predicate is factive if and only if its CC is presupposed and entailed, this means that these two predicates are factive if their CCs are presupposed. As discussed above, experiments 1a,b found that the CCs of all twenty predicates investigated were at least mildly projective, compared to main clause content. We are therefore led to the conclusion that, according to the definition in 3b, these two predicates are the best contenders for factive predicates from our set of twenty. This is a very unsatisfying conclusion, because the CCs of these two predicates were among the least projective.

The second unexpected implication concerns some of the predicates for which experiment 2a found their CCs not to be entailed, namely *see*, *discover*, *know*, *be annoyed*, *reveal*, and *demonstrate*. Under the definition of factive predicates in 3b, this result

²⁷ A Bayesian mixed-effects linear regression with the same fixed- and random-effects structure yielded similar results, except that in addition to *prove* and *be right*, there was no evidence for a difference between the entailing controls and *know*, *confirm*, *discover*, and *see*. See the Github repository mentioned in n. 14 for the model output.

means that, contrary to what is canonically assumed, the predicates *see*, *discover*, *know*, *be annoyed*, and *reveal* are not factive, nor is the predicate *demonstrate* veridical.

Our way of identifying entailed CCs, namely by comparing the CCs of the twenty predicates to entailed content, was motivated by the standard definition of entailment. As with the results of experiment 1a, one can, of course, ask whether there might be better ways of identifying entailed CCs. For instance, one might consider those CCs entailed whose mean inference rating was at least as high as that of *demonstrate* (0.85): this would mean that the CCs of all of the predicates typically taken to be factive or veridical non-factive are entailed, as well as the CCs of some optionally factive predicates. There is, however, no principled reason why the rating for *demonstrate* should be the threshold for entailment. One could also pick an arbitrary mean inference rating, say 0.9, and categorize CCs with mean ratings of at least 0.9 as entailed, and those with mean ratings below 0.9 as not entailed. Of course, this way of identifying entailed CCs is also not principled. After all, one could just as well pick a threshold of 0.95 or 0.85, with the result that the CCs of different predicates would be considered entailed: for instance, the CC of *reveal* (0.9) would count as entailed with a threshold of 0.85 but not with a threshold of 0.95. Furthermore, neither of these alternative ways of identifying entailed CCs is compatible with the standard definition of entailment, because, under both, one would need to allow CCs to count as entailed even though they received significantly lower ratings than entailed content.

It is also possible, however, that the surprising result of experiment 2a—that *see*, *discover*, *know*, *be annoyed*, and *reveal* are not factive under definition 3b—is an artifact of the response task. For instance, the gradient response scale on which participants responded in experiment 2a may have contributed to lower inference ratings for the CCs of these predicates, resulting in their classification as not entailed and, therefore, not factive under definition 3b. To assess this possibility, we conducted experiment 2b, which was identical to experiment 2a but employed a two-alternative forced-choice task.

3.2. EXPERIMENT 2B: CATEGORICAL RATINGS USING THE INFERENCE DIAGNOSTIC FOR ENTAILMENT.

METHODS.

Participants. Six hundred participants with US IP addresses and at least 99% of previous HITs approved were recruited on Amazon’s Mechanical Turk platform. They were paid \$1.

What is true: Brian is annoyed that Zoe calculated the tip.	
Does it follow that Zoe calculated the tip?	
<input type="radio"/> No	<input type="radio"/> Yes
Next	

FIGURE 10. A sample trial in experiment 2b.

Materials and procedure. The materials and procedure were identical to those of experiment 2a, with the exception of the task: participants responded ‘yes’ or ‘no’ to the question of whether the CC follows from the true statement, as shown in Figure 10.

Data exclusion. The data from 225 participants were excluded according to the criteria given in supplement §C. The data from 375 participants (ages: 18–73, median: 38; 187 female, 187 male, 1 undeclared) were analyzed.

RESULTS AND DISCUSSION. Figure 11 shows the proportion of ‘yes’ ratings by predicate, in increasing order from left to right, as well as for nonentailing and entailing controls. The results were strikingly similar to those of experiment 2a: the proportions of ‘yes’ responses were quite high for most factive predicates as well as for the veridical

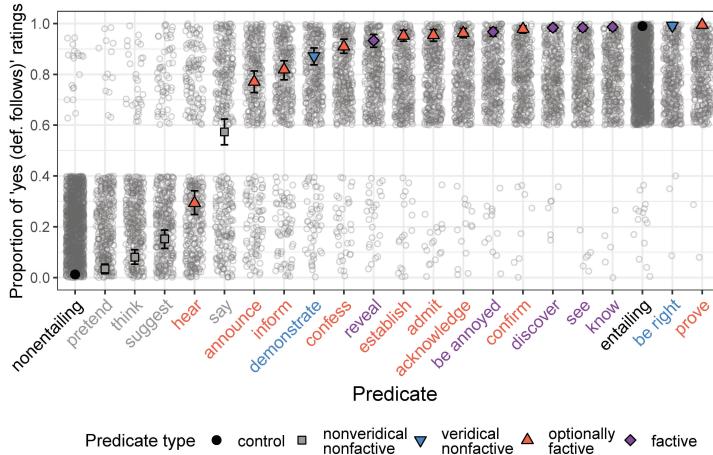


FIGURE 11. Proportion of ‘yes’ ratings by predicate in experiment 2b. Error bars indicate 95% bootstrapped confidence intervals. Vertically and horizontally jittered light gray dots indicate individual participants’ responses of ‘yes’ (coded as 1) and ‘no’ (coded as 0).

nonfactive predicate *be right*, but lower for the canonically factive predicate *reveal* and the canonically veridical nonfactive predicate *demonstrate*. The ordering of the predicates is almost entirely identical to the gradient measure of experiment 2a, as evidenced by the very high Spearman rank correlation of 0.989 between experiments 2a and 2b (see supplement §E for a visualization).

To assess which predicates have entailed CCs, we fitted a Bayesian mixed-effects logistic regression model with weakly informative priors using the brms package in R, which predicted the log-odds of a ‘yes’ over a ‘no’ response from a fixed effect of predicate (with treatment coding and ‘entailing control’ as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. The estimated log-odds of a ‘yes’ response was lower for each predicate than for the entailing controls: that is, the 95% credible intervals for the estimated coefficients did not contain 0 for any predicate, except for *prove*, *be right*, *know*, *see*, *discover*, and *confirm*. Thus, the categorical inference ratings for the CCs of *prove*, *be right*, *know*, *see*, *discover*, and *confirm* are compatible with the assumption that the CCs of these predicates are entailed. The categorical inference ratings for the remaining predicates, including the canonically factive predicates *be annoyed* and *reveal* and the veridical nonfactive predicate *demonstrate*, were significantly lower, suggesting that these CCs are not entailed, contrary to assumption.

Changing the response task from a gradient one (experiment 2a) to a categorical one (experiment 2b) resulted in slightly different conclusions: whereas experiment 2a supports only the assumption that the CCs of *be right* and *prove* are entailed, experiment 2b

additionally supports the assumption that the CCs of *know*, *see*, *discover*, and *confirm* are entailed. However, the implications of experiment 2b are just as unexpected as those of experiment 2a. First, because the CCs of *confirm*, *discover*, *see*, *know*, *be right*, and *prove* are entailed and projective, this set of predicates is identified, by the definition in 3b, as factive. This is an unsatisfying conclusion, because this set of predicates is very heterogeneous with respect to projection: the CC of *be right* was among the least projective of the twenty predicates investigated, while that of *know* was among the most projective. Thus, this set of predicates can hardly be considered a natural class. Second, the predicates *be annoyed* and *reveal* are not factive under the definition of factive predicates in 3b, and the predicate *demonstrate* is not veridical.

The results of both experiments 2a and 2b have unexpected implications for the identification of factive predicates. One might ask whether these results and their implications are an artifact of the inference diagnostic for entailment. To assess this possibility, we conducted experiments 3a and 3b, which were identical to experiments 2a and 2b, respectively, but employed the contradictoriness diagnostic for entailment.

3.3. EXPERIMENT 3A: GRADIENT RATINGS USING THE CONTRADICTORINESS DIAGNOSTIC FOR ENTAILMENT. Experiment 3a investigated which of the CCs of the twenty clause-embedding predicates are entailed, based on the contradictoriness diagnostic for entailment in 20b. Participants rated the contradictoriness of utterances of English sentences of the form ϕ *but not* ψ , where ϕ is an indicative matrix sentence with a clause-embedding predicate and ψ is its clausal complement, as in 25.

- (25) Mary announced that she is pregnant, but she's not.
Gradient contradictoriness ratings were collected.

METHODS.

Participants. Three hundred participants with US IP addresses and at least 99% of previous HITs approved were recruited on Amazon's Mechanical Turk platform. They were paid \$0.75.

Materials. On target trials, the 400 predicate + clause combinations were combined with a proper-name subject and a *but*-clause that denied the truth of the CC. As shown in 26, stimuli were presented to participants as utterances by named speakers. The proper names that realized the speakers, the subjects of the twenty predicates, and the subjects of the complement clauses were all unique. The gender of the proper-name subject of the predicate was distinct from the gender of the proper name in the complement clause, to ensure that the pronoun in the elliptical *but*-clause unambiguously referred to the individual in the complement clause.

- (26) a. **Christopher:** Melissa knows that Danny ate the last cupcake, but he didn't.
b. **Susan:** Jerry pretended that Emma studied on Saturday morning, but she didn't.

The experiment included eight control stimuli that were used to assess whether participants were attending to the task, as well as to identify contradictoriness ratings for entailed content. There were four control stimuli, like 27, for which we expected contradictoriness ratings to be at ceiling because the first clause of each of these sentences entails the negation of the second clause. In 27, for instance, the content of *Madison laughed loudly* entails that Madison laughed, that is, the negation of *she didn't laugh*. As discussed above, we relied on the contradictoriness ratings for these control stimuli to identify entailments. By contrast, we expected the contradictoriness ratings for the four noncontradictory control stimuli, like 28, to be at floor, because the content of the second clause does not contradict the content of the first clause. For instance, in 28,

the content that Vanessa is good at math does not entail the negation of the content that the speaker is not good at math. For the full set of control stimuli see supplement §B.

(27) Madison laughed loudly and she didn't laugh.

(28) Vanessa is really good at math, but I'm not.

Each participant saw a random set of twenty-eight stimuli: each set contained one target stimulus for each of the twenty predicates (each with a unique complement clause) and the same eight control stimuli. Trial order was randomized.

Margaret: "Edward heard that Mary is pregnant, but she isn't."

Is Margaret's utterance contradictory?

definitely no definitely yes

[Next](#)

FIGURE 12. A sample trial in experiment 3a.

Procedure. Participants were told that they would read utterances produced by a speaker and were asked to assess whether the utterance is contradictory. On each trial, participants read the speaker's utterance and then gave their response on a slider marked 'definitely no' at one end (coded as 0) and 'definitely yes' at the other (coded as 1), as shown in Figure 12. Participants first responded to two familiarization stimuli.

Data exclusion. The data from thirty-seven participants were excluded according to the criteria given in supplement §C. The data from 263 participants (ages: 18–72, median: 36; 126 female, 136 male, 1 other) were analyzed.

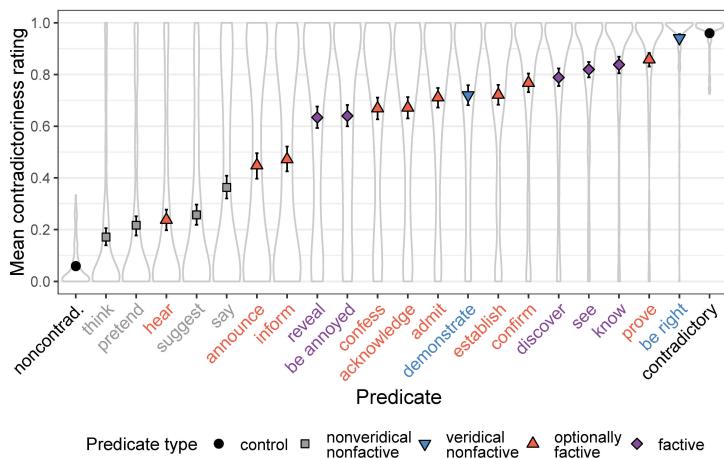


FIGURE 13. Mean contradictoriness ratings by predicate in experiment 3a, including the noncontradictory and contradictory controls. Error bars indicate bootstrapped 95% confidence intervals. Violin plots indicate the kernel probability density of the individual participants' ratings.

RESULTS AND DISCUSSION. Figure 13 shows the mean contradictoriness ratings for the target and control stimuli. In line with impressionistic judgments reported in the literature, the mean inference ratings for the canonically factive and veridical nonfactive predicates were among the highest, and those of the nonfactive nonveridical predicates were among the lowest. However, except possibly for *be right*, none of the mean contradictoriness ratings were close to those of the contradictory controls.

To assess which predicates have entailed CCs, we fitted a Bayesian mixed-effects beta regression model with weakly informative priors using the brms package in R, which predicted the contradictoriness ratings from a fixed effect of predicate (with treatment coding and ‘contradictory control’ as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. According to this model, the estimated mean for all predicates was lower than that of the contradictory controls: that is, the 95% credible intervals for the estimated adjustment to the intercept mean did not contain 0 for any predicate. See supplement §D for the full model output.²⁸ Thus, the gradient contradictoriness ratings for none of the predicates are compatible with the assumption that the CCs are entailed, and, consequently, none of the twenty predicates are factive under the definition in 3b.

Given the difference in results for experiments 2a and 2b, one may again ask whether the lack of entailed CCs identified by experiment 3a is an artifact of the gradient response task, and whether categorical responses might result in more CCs being classified as entailed. To assess this possibility, we conducted experiment 3b, which was identical to experiment 3a but employed a two-alternative forced-choice task.

3.4. EXPERIMENT 3B: CATEGORICAL RATINGS USING THE CONTRADICTORINESS DIAGNOSTIC FOR ENTAILMENT.

METHODS.

<p>Margaret: "Edward informed Sam that Sophia got a tattoo, but she didn't."</p> <p>Is Margaret's utterance contradictory?</p> <p style="text-align: center;"><input type="radio"/> No <input type="radio"/> Yes</p> <p style="text-align: center;">Next</p>

FIGURE 14. A sample trial in experiment 3b.

Participants. Six hundred participants with US IP addresses and at least 99% of previous HITs approved were recruited on Amazon’s Mechanical Turk platform. They were paid \$1.

Materials and procedure. The materials and procedure were identical to those of experiment 3a, with the exception of the task: participants responded ‘yes’ or ‘no’ to the question of whether the speaker’s utterance is contradictory, as shown in Figure 14.

Data exclusion. The data from 247 participants were excluded according to the criteria given in supplement §C. The data from 353 participants (ages: 18–73, median: 37; 180 female, 173 male) were analyzed.

²⁸ A Bayesian mixed-effects linear regression with the same fixed- and random-effects structure yielded qualitatively identical results, except that the contrast between *be right* and the contradictory controls was not significant. See the Github repository mentioned in n. 14 for the model output.

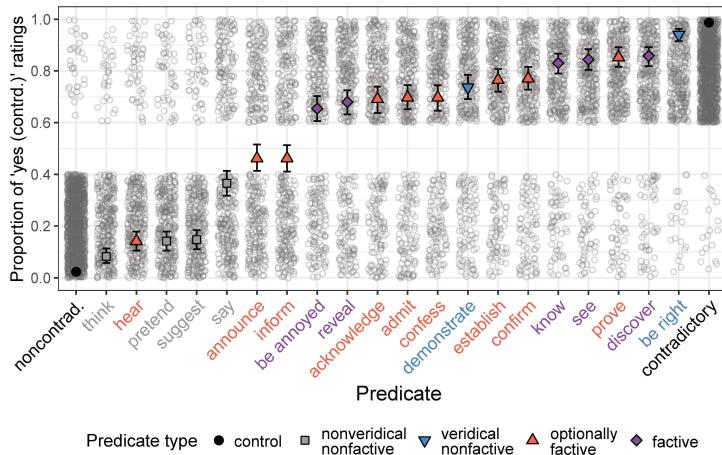


FIGURE 15. Proportion of ‘yes’ ratings by predicate in experiment 3b. Error bars indicate 95% bootstrapped confidence intervals. Vertically and horizontally jittered light gray dots indicate individual participants’ responses of ‘yes’ (coded as 1) and ‘no’ (coded as 0).

RESULTS AND DISCUSSION. Figure 15 shows the proportion of ‘yes’ responses on target and control trials. The results of experiment 3b were very similar to those of experiment 3a: except possibly for *be right*, the contradictoriness ratings for the CCs of none of the predicates were close to those of the contradictory controls. The Spearman rank correlation between experiments 3a and 3b was very high, at 0.985 (see supplement §E for a visualization).

To assess which predicates have entailed CCs, we fitted a Bayesian mixed-effects logistic regression model with weakly informative priors using the *brms* package in R, which predicted the log-odds of a ‘yes’ over a ‘no’ response from a fixed effect of predicate (with treatment coding and ‘contradictory control’ as reference level) and included the maximal random-effects structure justified by the design, namely random by-participant and by-item intercepts. The estimated log-odds of a ‘yes’ response was lower for each predicate than for the contradictory controls: that is, the 95% credible intervals for the estimated coefficients did not contain 0 for any predicate. Thus, as with experiment 3a, the results of experiment 3b suggest that the CCs of none of the twenty predicates are entailed. Consequently, as with experiment 3a, none of the twenty predicates are factive under the definition of factive predicates in 3b.

3.5. INTERIM SUMMARY AND CONVERGING EVIDENCE. According to definition 3b, a predicate is factive if and only if its CC is presupposed and entailed. In experiments 2 and 3, we compared the inference ratings (experiments 2) and the contradictoriness ratings (experiments 3) for the CCs to the ratings for control stimuli with entailed content. We found that the CCs of *prove* and *be right* are entailed based on the gradient and the

²⁹ The differences in the results of experiments 2 and 3 highlight the importance for future research of investigating the validity of entailment diagnostics. While the ratings were overall lower with the contradictoriness diagnostic than with the inference diagnostic, the high Spearman rank correlations between the ratings from the two diagnostics—0.954 for the two gradient measures and 0.934 for the two categorical measures—suggests that both are measuring a similar underlying concept. Moreover, the entailed controls received very high ratings across experiments 2 and 3, suggesting that all four measures are able to diagnose entailed content with some reliability. However, de Marneffe et al. (2012:329) have suggested that veridicality ratings (comparable to our inference ratings) are influenced by pragmatic reasoning; see also Pavlick & Kwiatkowski 2019.

categorical inference diagnostic (experiments 2a and 2b), that the CCs of *see*, *discover*, and *confirm* are entailed based on the categorical inference diagnostic (experiment 2b) but not on the gradient one (experiment 2a), and that the CCs of none of the twenty predicates are entailed based on the gradient and categorical contradictoriness diagnostics for entailment (experiments 3a,b).²⁹ Given that the CCs of all twenty predicates are at least mildly projective, this means, under definition 3b, that the predicates *prove* and *be right* are factive according to the results of experiment 2a, that the predicates *prove*, *be right*, *know*, *see*, *discover*, and *confirm* are factive according to the results of experiment 2b, and that none of the twenty predicates are factive according to the results of experiments 3a,b.

These results lead us to the conclusion that projection and entailment jointly do not identify a coherent class of factive predicates, contrary to what is expected on definition 3b. First, given the results of the gradient inference diagnostic (experiment 2a), the predicates *be right* and *prove* are factive, but it is unsatisfying to consider them factive predicates because they both have only weakly projective CCs. Second, given the results of the categorical inference diagnostic (experiment 2b), the predicates *prove*, *be right*, *see*, *discover*, *confirm*, and possibly *know* are factive, but this is a heterogeneous set of predicates and not a natural class: the CC of *be right* was among the least projective and that of *know* among the most projective. Finally, given the results of the contradictoriness diagnostic, none of the twenty predicates are factive. In short, the set of predicates identified as factive is either heterogeneous with respect to projection, that is, not a natural class, or empty.

Converging evidence for our conclusion comes from the VerbVeridicality and MegaVeridicality data sets. Consider first the VerbVeridicality data set: in addition to projection ratings for the CCs of the negated variants of the 859 indicative matrix sentences, it also contains entailment ratings for the 859 indicative matrix sentences; an example is given in 29.

- (29) The GAO has indicated that it is unwilling to compromise.

(Ross & Pavlick 2019:2234)

To diagnose entailment, three participants rated on a five-point Likert scale whether the CC is true given the target sentence; one endpoint of the scale was labeled ‘definitely true’ (coded as 2) and the other ‘definitely not true’ (coded as -2). We follow Ross and Pavlick (2019) in referring to these as VERIDICALITY RATINGS.

Figure 16 plots the mean veridicality ratings for the seventy-eight clause-embedding predicates in the VerbVeridicality data set, with labels for the fifteen predicates from our experiments (recall that *see* and *saw* are plotted separately). On the standard definition of entailment, predicates with entailed CCs should have a mean veridicality rating of 2. None of the clause-embedding predicates have a mean veridicality rating of 2: thus, the CCs of none of these seventy-eight clause-embedding predicates are entailed, and, consequently, none of the seventy-eight predicates are factive according to the definition in 3b. Relaxing the linking function results in a heterogeneous set of factive predicates: for instance, there are fifteen predicates whose CC received a mean veridicality rating at least as high as that of *know* (1.5). If one assumes that these CCs are en-

³⁰ The VerbVeridicality data set does not include projection ratings for main clause content to which the projection ratings for the CCs could be compared. The cutoff of a mean projection rating of 0.3 (on a scale from -2 to 2) is arbitrary but comparable to the mean certainty rating of 0.15 (on a scale from 0 to 1) for *pretend* in our experiment 1a.

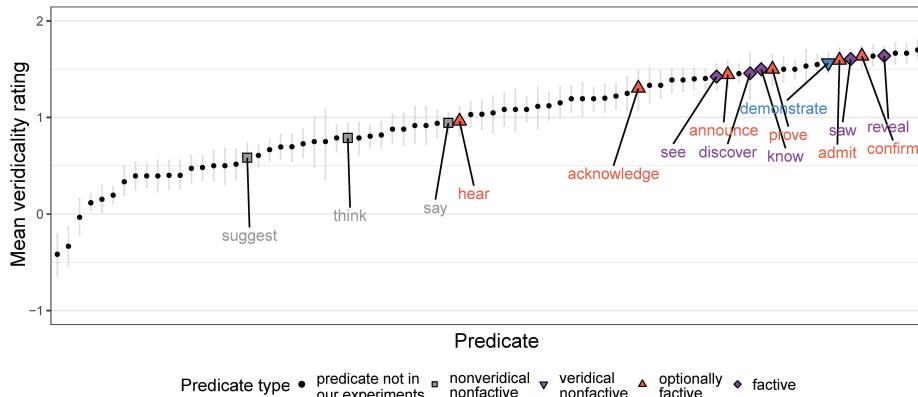


FIGURE 16. Mean veridicality ratings by predicate, with 95% bootstrapped confidence intervals, for the seventy-eight predicates in the VerbVeridicality data set (Ross & Pavlick 2019), with labels for the fifteen predicates featured in our experiments.

tailed and that CCs with mean projection ratings of at least 0.3 are presupposed,³⁰ then the set of factive predicates is heterogeneous with respect to projection: it includes *notice*, with a mean projection rating of 1.1, *give* (1.1), *realize* (1.1), *explain* (0.9), and *know* (0.8), but also *acknowledge* (0.5), *learn* (0.4), *saw* (0.4), and *admit* (0.3). In short, on the VerbVeridicality set, too, projection and entailment jointly do not identify a coherent class of factive predicates, contrary to what is expected on definition 3b.

The MegaVeridicality data set includes veridicality ratings for the CCs of the same 517 English clause-embedding predicates for which projection ratings were collected. The veridicality ratings were collected for stimuli that combined these predicates with arguments with low lexical content, as shown in 30 for *know*. To diagnose veridicality, participants were asked to respond to the question *Did that thing happen?*; the response options were ‘yes’ (which we again coded as 1), ‘maybe or maybe not’ (0), and ‘no’ (−1).

(30) Somebody knew that a particular thing happened.

Each of the 517 predicates in the MegaVeridicality data set received between nine and twenty entailment ratings (mean: ten) from 159 participants. Figure 17 plots the mean veridicality ratings, with labels for the nineteen predicates from our experiments. On the standard definition of entailment, one would expect predicates with entailed CCs to have a mean veridicality rating of 1. There are ninety-seven clause-embedding predicates that received only ‘yes’/1 ratings, including three of the canonically factive predicates we investigated, namely *reveal*, *know*, and *be annoyed*; this result is compatible with the assumption that the CCs of these ninety-seven predicates are entailed.³¹ If we assume, again, that CCs with mean projection ratings of at least 0.3 are presupposed, then ninety-two of the ninety-seven predicates are factive. This set is heterogeneous with respect to projection, including predicates like *bother*, with a mean projection rat-

³¹ It is not clear that the CCs of these ninety-seven predicates are entailed: this set includes *inform* (see the discussion in §1) as well as the communication predicates *bitch* and *howl*, whose CCs are not entailed. For instance, it does not follow indefeasibly from the naturally occurring example with *bitch* in (i) that the Democratic party hates all the Jews.

(i) Rambling to reporters in the Oval Office, Trump bitched that the Democratic party clearly hates all the Jews because of all the support reps Rashida Tlaib and Ilhan Omar received after he strong armed the Israelis into barring them from visiting the country as members of Congress.

(<https://www.wonkette.com/08-21-2019>)

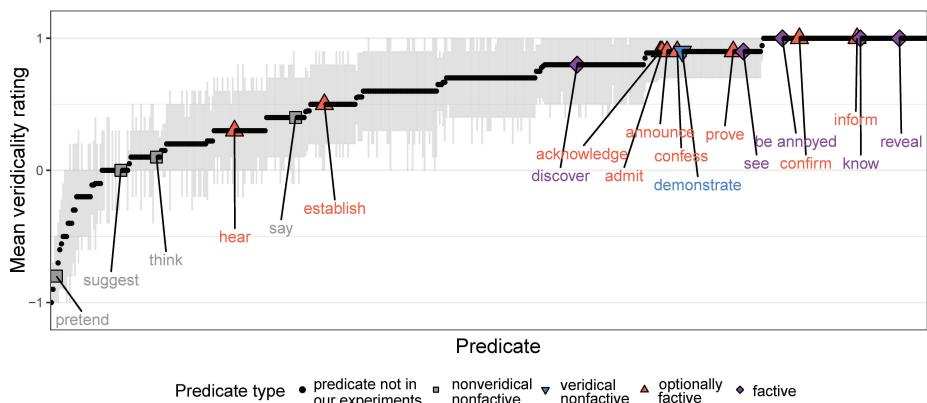


FIGURE 17. Mean veridicality ratings by predicate, with 95% bootstrapped confidence intervals, for the 517 predicates in the MegaVeridicality data set (White & Rawlins 2018, White et al. 2018), with labels for the nineteen predicates featured in our experiments.

ing of 0.97, *inform* (0.7), *know* (0.63), *grumble* (0.5), *notify* (0.4), *gloat* (0.4), *convey* (0.4), and *confirm* (0.3). Thus, on the MegaVeridicality set, too, projection and entailment jointly do not identify a coherent class of factive predicates.

This conclusion dovetails with that of White and Rawlins (2018), who assumed definition 3b. They pointed out (p. 228) that if one were to assume a majority response as threshold for whether the CC is presupposed and entailed, there are 199 factive and 292 nonfactive predicates in the MegaVeridicality data set, but also that ‘there are not necessarily clear dividing lines between these classes present in the data, suggesting that speaker’s inferences about [projection and entailment—*JD&JT*] are generally quite gradient and likely influenced by the fine-grained semantics of particular verbs’.

In sum, contrary to what is expected under the definition of factive predicates in 3b, the results of our experiments 1, 2, and 3, as well as those of the VerbVeridicality and MegaVeridicality data sets, suggest that projection and entailment of the CC do not jointly identify a class of factive predicates. Depending on how entailment is diagnosed and which linking function is assumed, the set of predicates identified as factive is either empty or heterogeneous with respect to projection.

4. GENERAL DISCUSSION. Sections 2 and 3 reported the results of experiments designed to investigate whether two prevalent definitions, repeated below, identify a class of factive predicates.

- (3) Two definitions of factive predicates: A clause-embedding predicate is factive if and only if the content of its clausal complement is
 - a. presupposed. (e.g. Kiparsky & Kiparsky 1970, Karttunen 1971a,b)
 - b. presupposed and entailed. (e.g. Gazdar 1979a, Schlenker 2010, Abrusán 2011, Anand & Hacquard 2014, Spector & Egré 2015)

The experiments in §2 investigated the definition in 3a: assuming projection as a diagnostic for presupposed CCs, the results of these experiments as well as the meta-analyses of three existing data sets suggest that projection does not identify a class of factive predicates. The experiments in §3 investigated the definition in 3b: the evidence from our experiments and meta-analyses of existing data sets suggests that, depending on how entailment is diagnosed, either there are no factive predicates or the set of predicates identified as factive is heterogeneous with respect to projection, that is, not a nat-

ural class. In sum, the results of our experiments do not provide empirical evidence for a coherent class of factive predicates under either of the two definitions in 3.

Our main motivation in conducting these investigations was to understand which predicates analyses of the projection of the CC of clause-embedding predicates apply to, as these analyses are generally limited to the CCs of factive predicates. The explanandum is assumed to be that the CCs of factive predicates, but not those of optionally factive or nonfactive ones, are presuppositions, which may project. On some analyses (e.g. Heim 1983, van der Sandt 1992), factive predicates lexically specify that the CC is required to be entailed by or satisfied in the relevant context; optionally factive and nonfactive predicates do not carry such a lexical specification. On other analyses (e.g. Abrusán 2011, 2016, Romoli 2015, Simons et al. 2017), projection of the CC is derived from the CC being entailed content that is backgrounded; nonentailed CCs of optionally factive and nonfactive predicates are again outside the scope of these analyses. Given the results of our experiments, we are unfortunately no closer to answering the question of which predicates these analyses apply to.³²

Where do we go from here? There are three broad ways forward. First, one might ask whether there are more appropriate diagnostics for presupposed or entailed content than the ones we have applied, or whether there are more suitable definitions of factive predicates than those in 3. Second, one might challenge our conclusion that the results of our investigations do not identify a class of factive predicates. We engage with several possible objections in §4.1. Alternatively, one could give up on the assumption that there is a coherent class of factive predicates (see also Karttunen 2016 and Djärv 2019 for such moves). In §4.2, we discuss what this move would mean for projection analyses.

4.1. POSSIBLE OBJECTIONS. The results of experiments 1a,b as well as the meta-analyses of the existing data sets in §2 are strongly suggestive of gradience in the projection of the CCs of clause-embedding predicates. Our discussions in §§2 and 3 suggested that this gradience precluded identification of a class of factive predicates on both definitions in 3. In the following, we consider several empirical, statistical, and conceptual objections to this conclusion.

OBJECTION 1: NOISE MIGHT HAVE OBSCURED AN UNDERLYINGLY BINARY FACTIVITY CATEGORY. Our results in experiments 1 are contingent on a particular operationalization of projection—via the ‘certain that’ diagnostic—and a particular behavioral measurement—via continuous slider ratings and binary ‘yes’/‘no’ judgments. There are multiple ways in which experimental noise might have been introduced and might have masked an underlying binary factivity category: (i) via the experimental items, which did not consist of the predicates alone, but of sentences in which the predicates were combined with particular subject noun phrases and complement clauses, (ii) and which moreover were produced by a named speaker, (iii) via the response tasks, and (iv) via

³² To be clear: contemporary analyses of the projection of the CCs of factive predicates can predict the non-projection of factive CCs in particular utterances. For instance, in Heim 1983 and van der Sandt 1992, the CC is locally accommodated when its projection would result in a contradictory or uninformative utterance. Although projection of the CC does not result in contradictory or uninformative utterances in the stimuli in our experiments 1a and 1b, one might assume, as suggested by an anonymous referee, (i) that some participants enrich the context in which the polar-question stimuli were presented with information that does result in a contradiction or un informativity, and (ii) that participants are more likely to enrich the context with such information with some factive predicates than others. Under these assumptions, these analyses may be compatible with the results of our experiments 1a,b, namely, that the projection of the CC of canonically factive predicates is heterogeneous and not categorically higher than that of other predicates. Neither assumption, however, strikes us as independently motivated.

participant-internal noise. Each of (i) to (iv) could have introduced systematic noise into the ratings, yielding gradient responses for predicates that may underlyingly be either factive or nonfactive.

The experiments reported here were designed to mitigate several possible sources of systematic noise: (i) all predicates occurred in the same sentence structures and were combined with the same set of complement clauses across participants, (ii) all predicates were presented in utterances situated in the same context, (iii) ratings were collected across different tasks, and (iv) ratings were obtained from a large number of participants, whose response variability was taken into consideration in the statistical analysis. Reasons to believe that these decisions mitigated noise include the replication in gradient orderings of predicates across continuous (slider ratings) and categorical (selections) response measures (experiment 1a vs. 1b); the replication of experiment 1a in experiment 1 of Degen & Tonhauser 2021 with a Spearman rank correlation of 0.991; and the replication of gradient orderings of predicates in the meta-analyses of existing data sets, which varied in the experimental tasks and items. We therefore consider it unlikely that experimental noise might have masked an underlying binary factivity category, though there is much need for discussion in the experimental semantics/pragmatics literature about the relation between behavioral data and semantic/pragmatic categories.³³

OBJECTION 2: THE DISTINCTION BETWEEN FACTIVE AND NONFACTIVE PREDICATES MIGHT BE FUZZY. Another objection to the conclusion that projection does not identify a class of factive predicates is that the expectation that the CCs of canonically factive predicates should be more projective than those of all other predicates is too strong. Specifically, given that the CCs of canonically factive predicates are among the most projective in both experiments 1a and 1b, one might argue that to motivate a class of factive predicates it suffices to show that assuming such a class explains variability observed in the projection of the CCs of clause-embedding predicates. To engage with this possibility, we analyzed the certainty ratings of experiments 1a,b in a Bayesian mixed-effects beta regression model (experiment 1a) and a logistic regression model (experiment 1b) that predicted projection from a binary factivity predictor, whereby the five predicates in 13a were classified as factive and the remaining predicates (as well as the main clause controls) as nonfactive. We report mean estimates and 95% credible intervals for the effect of this binary factivity predictor. Under such an analysis of the predicates, the CCs of presumed factive predicates were indeed more projective, both in experiment 1a ($\beta = 1.49$, 95% CIs [1.30, 1.67]) and in experiment 1b ($\beta = 3.99$, 95% CIs = [3.40, 4.59]).

A way to ask statistically whether binary factivity is a useful category to posit is to ask how much variability in the data is captured by the binary factivity models compared to the models that include the individual predicates as predictors, which were reported in §2.1 and §2.2. We refer to these latter models as predicate-specific models. Model comparison revealed that predicate-specific models were favored over binary factivity models in both experiment 1a (WAIC of predicate-specific model: -18971.3 ($SE = 359.3$); WAIC of binary factivity model: -18643.9 ($SE = 355.4$)) and experiment

³³ For such a discussion in the experimental syntax literature, see Keller 2000, Sorace & Keller 2005, Sprouse 2007, Hofmeister & Sag 2010, Sprouse & Almeida 2013, Lau et al. 2014, 2017.

³⁴ The WIDELY APPLICABLE INFORMATION CRITERION (WAIC) is a measure of model quality that estimates the pointwise out-of-sample prediction accuracy from a fitted Bayesian model using the log-likelihood evaluated at the posterior simulations of the parameter values and penalizes models with more parameters over models with fewer (Watanabe 2010). We obtained WAIC values for each model using the ‘loo’ package (Vehrtari et al. 2017). Lower WAIC values indicate better models.

1b (WAIC of predicate-specific model: 6444.5 ($SE = 124.8$); WAIC of binary factivity model: 6719.0 ($SE = 123.7$)).³⁴ This suggests that the qualitative observation made in §2.1 and §2.2, namely, that projection variability is better captured by taking into account the individual predicates than a binary factivity classification, is borne out in rigorous model comparison.

OBJECTION 3: THERE MAY BE UNCERTAINTY ABOUT WHICH CATEGORY A PREDICATE BELONGS TO. It has been suggested to us, in personal communication and by a referee, that the observed gradience in projection may be compatible with a binary factivity category in combination with two assumptions: first, that predicates may be ambiguous between a factive lexical entry, on which the CC is presupposed and hence projects, and a nonfactive lexical entry, on which the CC is not presupposed (for such a proposal for clause-embedding predicates see Spector & Egré 2015:1736; for a similar proposal for evaluative adjectives see Karttunen et al. 2014); and, second, that interpreters may be uncertain about which lexical entry a speaker intended in their utterance. Under such a view, a categorical notion of factivity is upheld, and the observed projection gradience can be hypothesized to be due to listeners' uncertainty about the use of any given predicate, which may be resolved preferentially one way or another through pragmatic factors.

The suggestion is that one way to test whether it is possible to maintain a binary factivity category with these two assumptions is to fit so-called finite mixture models to the data (Jacobs et al. 1991, Richardson & Green 1997). These models allow for representing the possibility that multiple components, that is, multiple subsets, make up an overall data set. For instance, the data for each predicate from experiment 1a could be analyzed with mixture models with respect to how many components (one, two, three, four, five, ...) best account for the data. To illustrate, consider Figure 18, which shows simulated distributions of certainty ratings together with the optimal number of Gaussian components.³⁵ Given the aforementioned assumptions, one could attribute certainty ratings that are best captured with one Gaussian component with a mean of 0.4, as in panel (a), to a predicate that only has a nonfactive lexical entry. Likewise, one could attribute certainty ratings that are best captured with one Gaussian component with a mean of 0.85, as in panel (b), to a predicate that only has a factive lexical entry. Certainty ratings like those in panels (c) and (d), which are best captured by two Gaussian components, might then be attributed to predicates that have both a factive and a nonfactive lexical entry, with differences between (c) and (d) being due to the predicates interacting differently with pragmatic factors that modulate projection. The projection gradience observed in experiment 1a could be hypothesized to come about by virtue of one or two components being combined in various ways.

We caution against this approach on conceptual grounds. What mixture models can do is supply us with an answer to the question of how many components optimally characterize the data. Mixture models are not, however, a theory of factivity or projection, and as such they do not specify what the generative model underlying the components is. Thus, while mixture models can be used to test the predictions of theories of factivity or projection that provide us with such a generative model, we caution against their application in the absence of such a theory. And, in fact, we do not currently have such a theory. The account sketched above, on which predicates may have a factive lexical entry, a nonfactive one, or both, does not lead us to expect projection data distribu-

³⁵ Models were fit using the ‘mclust’ and ‘mixtools’ packages (Benaglia et al. 2009, Scrucca et al. 2016) in R.

³⁶ An additional problem: for simplicity we introduced mixture models with Gaussian mixtures. A preliminary investigation of the data from experiment 1a via mixture models suggests that fifteen of the twenty pre-

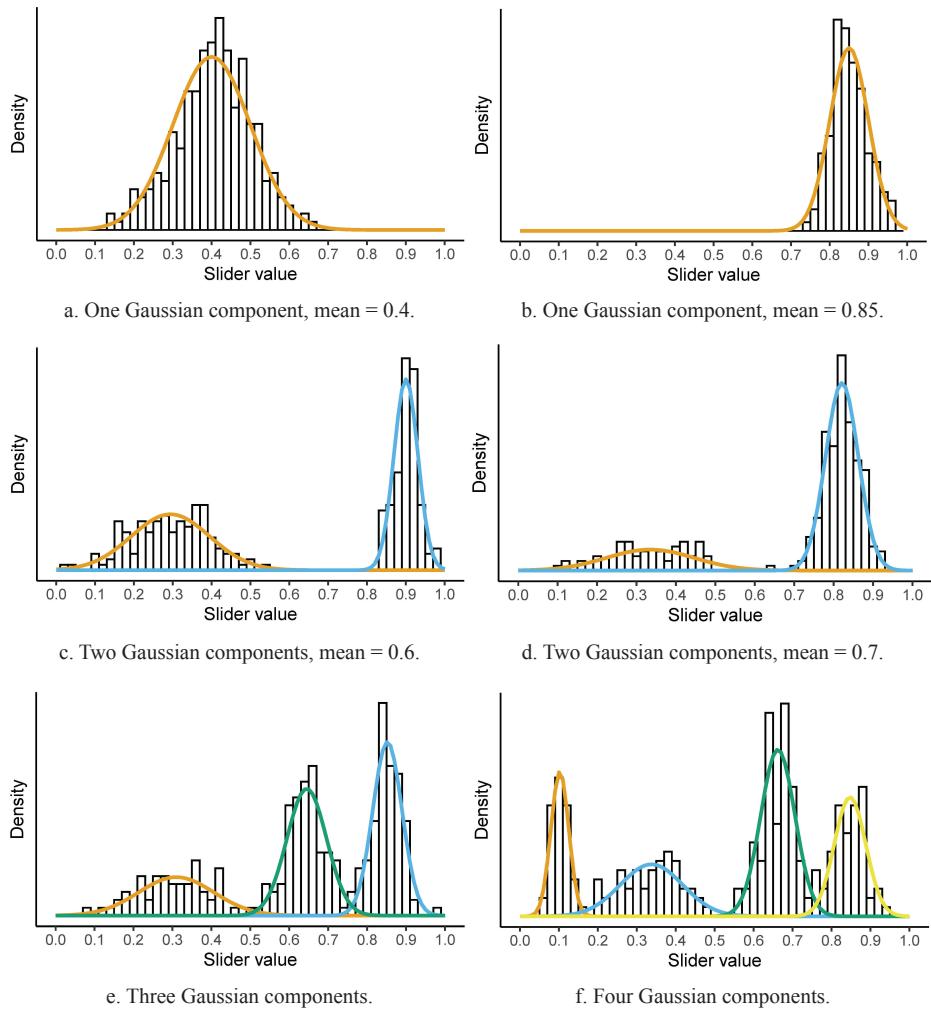


FIGURE 18. Different simulated ratings distributions with overlaid optimal number of Gaussian components.

tions, like those in panels (e) and (f) of Fig. 18, that are best captured by three or four components, let alone distributions that are best captured by five or six components. Worse yet, in the absence of such a theory, we would not know how to interpret data that are best captured by more than two components.³⁶ In short, we currently lack a principled reason to apply mixture models to the data reported here.

SUMMARY. We have argued that the results of our experiments and the meta-analyses of existing data sets did not find empirical support for a class of factive predicates. Our discussion in §4.1 engaged with several objections to this conclusion. In doing so, the discussion highlighted several empirical, statistical, and conceptual considerations for future research on the classification of clause-embedding predicates.

icates are best characterized by three to six components; none are best characterized by just one component (see supplement §F). However, the assumption of normality made by Gaussian mixture models does not hold in the bounded slider ratings data of experiment 1a. One would need to fit mixtures of beta distributions, for which statistical inference about optimal cluster number is much more complex.

4.2. IMPLICATIONS FOR PROJECTION ANALYSES. The assumption that there is a class of factive clause-embedding predicates has played a central role in linguistic theorizing on a broad range of topics (see the list in §1). In this section, we discuss the implications of our findings for one of these topics, namely projection analyses. As mentioned above, the explanandum of such analyses is assumed to be that the CCs of factive predicates, but not those of optionally factive or nonfactive ones, are presuppositions, which may project. If future research provides empirical support for a class of factive predicates (see §4.1), then it is the predicates in that class that these analyses may apply to.

But even in the absence of a conclusive list of factive predicates, the results of our experiments 1a,b as well as the meta-analyses of the three existing data sets point to exciting avenues for future research on projective content. First, our results show that the CCs of many more clause-embedding predicates than previously assumed are projective. This suggests that research on projective content has a much broader empirical scope than previously assumed. Given that research on projective content has focused almost entirely on the CC of canonically factive predicates, to the exclusion of the CC of other clause-embedding predicates, our results mean that we are now tasked with investigating and formally analyzing the projection of the CC of a much broader set of clause-embedding predicates than before. Research on projective content has identified several factors that influence the projection of utterance content, including context, lexical content, prior beliefs, at-issueness, information structure, and the perceived degree of reliability and trustworthiness of the attitude holder (e.g. Gazdar 1979a,b, Beaver 2010, Schlenker 2010, Simons et al. 2010, Abrusán 2011, 2016, Anand & Hacquard 2014, Cummins & Rohde 2015, Tonhauser 2016, 2020, Djärv & Bacovcin 2017, Simons et al. 2017, Tonhauser et al. 2018, Tonhauser et al. 2019, Mahler 2020, Mahler et al. 2020, Degen & Tonhauser 2021). An exciting question for future research is whether these factors are implicated in the projection of the CCs of all clause-embedding predicates, or only of particular subsets, and how they contribute to the observed projection variability between clause-embedding predicates. A framework that offers a principled mechanism for modeling the probabilistic interaction of factors like lexical content, speaker reliability, and prior beliefs about likely meanings is the RATIONAL SPEECH ACT (RSA) framework (Bergen et al. 2016). We believe that implementing the interpretation of utterances with clause-embedding predicates in RSA offers a promising way forward (for RSA models of other projective contents see Qing et al. 2016 and Stevens et al. 2017).

Second, our results show that there is a lot of variability in how projective the CCs of the different clause-embedding predicates are. One of the factors that clearly plays a role in predicting projection variability is the lexical meaning of the clause-embedding predicate. Our experiments 1a and 1b, as well as the results presented in Tonhauser et al. 2018, de Marneffe et al. 2019, and Degen & Tonhauser 2021, suggest that a binary distinction between factive and nonfactive predicates is too coarse to adequately predict projection. We hypothesize that more fine-grained distinctions that are based on the lexical meaning and discourse use of clause-embedding predicates are critical; specifically, we hypothesize that the CCs of classes of predicates that share lexical meanings and discourse uses are similar in projection.³⁷

We are aware of two tentative proposals along these lines in the recent literature. Anand and Hacquard (2014) hypothesized that listeners may take speakers to be com-

³⁷ Similarly, White 2021 recently proposed that lexical semantic properties other than factivity may be implicated in another domain where factivity has been taken to play a role, namely, whether a predicate is responsive.

mitted to the CC of assertive predicates, like *acknowledge*, *admit*, and *confirm*, ‘because of the kind of discourse moves that these predicates report’ (p. 74). Specifically, utterances of sentences with these predicates can be used to report the success of an uptake in a reported common ground. For example, an utterance of a sentence like *Kim confirmed that Mary is the murderer* can report that the content *p* of the complement, that Mary is the murderer, is accepted ‘into the common ground of the reported discourse’ (*ibid.*). The authors suggested that ‘[t]his acceptance of *p* can easily bleed into the actual common ground, under the assumption that no subsequent move removed *p* from the common ground’ (pp. 74f.). Karttunen (2016) proposed different paths to projection for different classes of clause-embedding predicates. For instance, for predicates with *that*-clause subjects that do not involve an attitude holder, like *be odd*, *be tragic*, or *count*, the projection of the CC is derived from its status as a presupposition; with predicates that express a propositional attitude, like *know* or *forget*, the speaker is only necessarily committed to the attitude holder being committed to the truth of the CC, and speaker commitment to the CC is derived as a generalized conversational implicature; for a third subclass of change-of-state predicates, including *discover* and *notice*, the CC is entailed but its projection is derived pragmatically; and for a fourth subclass of communicative predicates, like *acknowledge*, *admit*, and *confess*, the speaker is merely committed to the attitude holder having communicated something that the attitude holder wishes to present as a fact. While such proposals need to be fleshed out in more detail to be predictive, we believe that detailed analyses of the lexical meanings of clause-embedding predicates are a fruitful path to a better understanding of the projection of their CCs.

5. CONCLUSION. Linguistic research since Kiparsky & Kiparsky 1970 has assumed that properties of the content of the clausal complement distinguish factive predicates from others. This article investigated whether the content of the complement of twenty English clause-embedding predicates is projective and entailed, with the goal of identifying whether there is a class of factive predicates according to two definitions found in the literature. Our experiments revealed that projection alone does not categorically distinguish canonically factive predicates from others, and that a coherent class of factive predicates also does not emerge from the additional consideration of entailment, as assessed by two standard diagnostics. In addition to considering how the search for a class of factive predicates might continue, we also suggested that our results open up several exciting new avenues for empirical and theoretical research on the projection of the contents of the complements of clause-embedding predicates.

REFERENCES

- ABRUSÁN, MÁRTA. 2011. Predicting the presuppositions of soft triggers. *Linguistics and Philosophy* 34.491–535. DOI: 10.1007/s10988-012-9108-y.
- ABRUSÁN, MÁRTA. 2014. *Weak island semantics*. New York: Oxford University Press. DOI: 10.1093/acprof:oso/9780199639380.001.0001.
- ABRUSÁN, MÁRTA. 2016. Presupposition cancellation: Explaining the ‘soft-hard’ trigger distinction. *Natural Language Semantics* 24.165–202. DOI: 10.1007/s11050-016-9122-7.
- ABUSCH, DORIT. 2010. Presupposition triggering from alternatives. *Journal of Semantics* 27(1).37–80. DOI: 10.1093/jos/ffp009.
- ANAND, PRANAV; JANE GRIMSHAW; and VALENTINE HACQUARD. 2019. Sentence embedding predicates, factivity and subjects. *Tokens of meaning: Papers in honor of Lauri Karttunen*, ed. by Cleo Condoravdi and Tracy Holloway King, 279–308. Stanford, CA: CSLI Publications.

- ANAND, PRANAV, and VALENTINE HACQUARD. 2014. Factivity, belief and discourse. *MIT Working Papers in Linguistics (The art and craft of semantics: A festschrift for Irene Heim)*, ed. by Luka Crnić and Uli Sauerland) 70.69–90. Online: <https://semanticsarchive.net/Archive/jZiNmM4N/>.
- ANDERSON, LLOYD B. 1986. Evidentials, paths of change, and mental maps: Typologically regular asymmetries. *Evidentiality: The linguistic coding of epistemology*, ed. by Wallace L. Chafe and Johanna Nichols, 273–312. Norwood, NJ: Ablex.
- ARAVIND, ATHULYA, and MARTIN HACKL. 2017. Factivity and at-issueness in the acquisition of *forget* and *remember*. *Proceedings of the Boston University Conference on Language Development (BUCLD)* 41.46–59.
- BEAVER, DAVID. 2001. *Presupposition and assertion in dynamic semantics*. Stanford, CA: CSLI Publications.
- BEAVER, DAVID. 2010. Have you noticed that your belly button lint color is related to the color of your clothing? *Presuppositions and discourse: Essays offered to Hans Kamp*, ed. by Rainer Bäuerle, Uwe Reyle, and Thomas Ede Zimmermann, 65–99. Leiden: Brill. DOI: 10.1163/9789004253162_004.
- BEAVER, DAVID, and BART GEURTS. 2014. Presupposition. *The Stanford encyclopedia of philosophy*, ed. by Edward N. Zalta. Stanford, CA: Stanford University. Online: <https://plato.stanford.edu/archives/win2014/entries/presupposition/>.
- BENAGLIA, TATIANA; DIDIER CHAUVEAU; DAVID R. HUNTER; and DEREK S. YOUNG. 2009. mixtools: An R package for analyzing mixture models. *Journal of Statistical Software* 32(6).1–29. DOI: 10.18637/jss.v032.i06.
- BERGEN, LEON; ROGER LEVY; and NOAH GOODMAN. 2016. Pragmatic reasoning through semantic inference. *Semantics and Pragmatics* 9:20. DOI: 10.3765/sp.9.20.
- BÜRKNER, PAUL-CHRISTIAN. 2017. brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software* 80(1).1–28. DOI: 10.18637/jss.v080.i01.
- CHIERCHIA, GENNARO, and SALLY McCONNELL-GINET. 1990. *Meaning and grammar*. Cambridge, MA: MIT Press.
- CREMERS, ALEXANDRE. 2018. Plurality effects in an exhaustification-based theory of embedded questions. *Natural Language Semantics* 26.193–251. DOI: 10.1007/s11050-018-9145-3.
- CUMMINS, CHRIS, and HANNAH ROHDE. 2015. Evoking context with contrastive stress: Effects on pragmatic enrichment. *Frontiers in Psychology* 6:1779. DOI: 10.3389/fpsyg.2015.01779.
- DE MARNEFFE, MARIE-CATHERINE; CHRISTOPHER D. MANNING; and CHRISTOPHER POTTS. 2012. Did it happen? The pragmatic complexity of veridicality assessment. *Computational Linguistics* 38(2).301–33. DOI: 10.1162/COLI_a_00097.
- DE MARNEFFE, MARIE-CATHERINE; MANDY SIMONS; and JUDITH TONHAUSER. 2019. The CommitmentBank: Investigating projection in naturally occurring discourse. *Proceedings of Sinn und Bedeutung* 23(2).107–24. DOI: 10.18148/sub/2019.v23i2.601.
- DE VILLIERS, JILL G. 2005. Can language acquisition give children a point of view? *Why language acquisition matters for theory of mind*, ed. by Janet Wilde Astington and Jodie A. Baird, 186–219. Oxford: Oxford University Press. DOI: 10.1093/acprof:oso/9780195159912.003.0010.
- DEGEN, JUDITH, and JUDITH TONHAUSER. 2021. Prior beliefs modulate projection. *Open Mind* 5.59–70. DOI: 10.1162/opmi_a_00042.
- DJÄRV, KAJSA. 2019. *Factive and assertive attitude reports*. Philadelphia: University of Pennsylvania dissertation. Online: <https://repository.upenn.edu/edissertations/3645>.
- DJÄRV, KAJSA, and HEZEKIAH AKIVA BACOVGIN. 2017. Prosodic effects on factive presupposition projection. *Proceedings of Semantics and Linguistic Theory (SALT)* 27.116–33. DOI: 10.3765/salt.v27i0.4134.
- DJÄRV, KAJSA; JÉRÉMY ZEHR; and FLORIAN SCHWARZ. 2016. Cognitive vs. emotive factives: An experimental differentiation. *Proceedings of Sinn und Bedeutung* 21(1).367–85. Online: <https://ojs.ub.uni-konstanz.de/sub/index.php/sub/article/view/143>.
- EGRÉ, PAUL. 2008. Question-embedding and factivity. *Grazer Philosophische Studien* 77. 85–125. DOI: 10.1163/18756735-90000845.
- GAZDAR, GERALD. 1979a. *Pragmatics: Implicature, presupposition and logical form*. New York: Academic Press.

- GAZDAR, GERALD. 1979b. A solution to the projection problem. *Syntax and semantics, vol. 11: Presupposition*, ed. by Choon-Kyu Oh and David Dinneen, 57–89. New York: Academic Press.
- GIANNAKIDOU, ANASTASIA. 1998. *Polarity sensitivity as (non)veridical dependency*. Amsterdam: John Benjamins. DOI: 10.1075/la.23.
- GIANNAKIDOU, ANASTASIA. 2001. The meaning of free choice. *Linguistics and Philosophy* 24.659–735. DOI: 10.1023/A:1012758115458.
- GIANNAKIDOU, ANASTASIA, and ALDA MARI. 2015. Mixed (non)veridicality and mood choice with emotive verbs. *Chicago Linguistic Society* 51.181–95.
- GIVÓN, TALMY. 1995. *Functionalism and grammar*. Amsterdam: John Benjamins. DOI: 10.1075/z.74.
- GUERZONI, ELENA, and YAEL SHARVIT. 2007. A question of strength: On NPIs in interrogative clauses. *Linguistics and Philosophy* 30.361–91. DOI: 10.1007/s10988-007-9014-x.
- HAZLETT, ALLAN. 2010. The myth of factive verbs. *Philosophy and Phenomenological Research* 80(3).497–522. DOI: 10.1111/j.1933-1592.2010.00338.x.
- HEIM, IRENE. 1983. On the projection problem for presuppositions. *West Coast Conference on Formal Linguistics (WCCFL)* 2.114–25.
- HEIM, IRENE. 1992. Presupposition projection and the semantics of attitude verbs. *Journal of Semantics* 9(3).183–221. DOI: 10.1093/jos/9.3.183.
- HEYCOCK, CAROLINE. 2006. Embedded root phenomena. *The Blackwell companion to syntax*, ed. by Martin Everaert and Henk van Riemsdijk, 174–209. Oxford: Blackwell.
- HINTIKKA, JAAKKO. 1975. Different constructions in terms of the basic epistemological verbs: A survey of some problems and proposals. *The intentions of intentionality and other new models for modalities*, ed. by Jaakko Hintikka, 1–25. Dordrecht: Reidel.
- HOFMEISTER, PHILIP, and IVAN A SAG. 2010. Cognitive constraints and island effects. *Language* 86(2).366–415. DOI: 10.1353/lan.0.0223.
- HUKARI, THOMAS E., and ROBERT D. LEVINE. 1995. Adjunct extraction. *Journal of Linguistics* 31(2).195–226. DOI: 10.1017/S002226700015590.
- JACOBS, ROBERT A.; MICHAEL I. JORDAN; STEVEN J. NOWLAN; and GEOFFREY E. HINTON. 1991. Adaptive mixtures of local experts. *Neural Computation* 3(1).79–87. DOI: 10.1162/neco.1991.3.1.79.
- KARTTUNEN, LAURI. 1971a. Implicative verbs. *Language* 47(2).340–58. DOI: 10.2307/412084.
- KARTTUNEN, LAURI. 1971b. Some observations on factivity. *Papers in Linguistics* 4.55–69. DOI: 10.1080/08351817109370248.
- KARTTUNEN, LAURI. 2016. Presupposition: What went wrong? *Proceedings of Semantics and Linguistic Theory (SALT)* 26.705–31. DOI: 10.3765/salt.v26i0.3954.
- KARTTUNEN, LAURI, and STANLEY PETERS. 1979. Conventional implicature. *Syntax and semantics, vol. 11: Presuppositions*, ed. by Choon-Kyu Oh and David A. Dinneen, 1–56. New York: Academic Press.
- KARTTUNEN, LAURI; STANLEY PETERS; ANNIE ZAENEN; and CLEO CONDORAVDI. 2014. The chameleon-like nature of evaluative adjectives. *Empirical Issues in Syntax and Semantics* 10.233–50. Online: http://www.cssp.cnrs.fr/eiss10/eiss10_karttunen-and-peters-and-zaenen-and-condoravdi.pdf.
- KARTTUNEN, LAURI, and ANNIE ZAENEN. 2005. Veridicity. *Annotating, extracting and reasoning about time and events* (Dagstuhl seminar proceedings 05151), ed. by Frank Schilder, Graham Katz, and James Pustejovsky. DOI: 10.4230/DagSemProc.05151.13.
- KASTNER, ITAMAR. 2015. Factivity mirrors interpretation: The selectional requirements of presuppositional verbs. *Lingua* 164(Part A).156–88. DOI: 10.1016/j.lingua.2015.06.004.
- KELLER, FRANK. 2000. *Gradience in grammar: Experimental and computational aspects of degrees of grammaticality*. Edinburgh: University of Edinburgh dissertation.
- KIPARSKY, PAUL, and CAROL KIPARSKY. 1970. Fact. *Progress in linguistics*, ed. by Manfred Bierwisch and Karl Erich Heidolph, 143–73. The Hague: Mouton. DOI: 10.1515/9783111350219.143.
- KLEIN, EWAN. 1975. Two sorts of factive predicates. *Pragmatics Microfiche* it 1.1. frames B5–C14.
- LANDAU, IDAN. 2001. *Elements of control: Structure and meaning in infinitival constructions*. Dordrecht: Springer. DOI: 10.1007/978-94-011-3943-4.

- LANGENDOEN, D. TERENCE, and HARRIS SAVIN. 1971. The projection problem for presuppositions. *Studies in linguistic semantics*, ed. by Charles J. Fillmore and D. Terence Langendoen, 54–60. New York: Holt, Rinehart and Winston.
- LASERSON, PETER. 2009. Relative truth, speaker commitment, and control of implicit arguments. *Synthese* 166.359–74. DOI: 10.1007/s11229-007-9280-8.
- LAU, JEY HAN; ALEXANDER CLARK; and SHALOM LAPPIN. 2014. Measuring gradience in speaker's grammaticality judgements. *Proceedings of the 36th annual meeting of the Cognitive Science Society (CogSci 2014)*, 821–26. Online: <https://cogsci.mindmodeling.org/2014/papers/149/>.
- LAU, JEY HAN; ALEXANDER CLARK; and SHALOM LAPPIN. 2017. Grammaticality, acceptability, and probability: A probabilistic view of linguistic knowledge. *Cognitive Science* 41(5).1202–41. DOI: 10.1111/cogs.12414.
- LORSON, ALEXANDRA. 2018. The influence of world knowledge on projectivity. Potsdam: University of Potsdam master's thesis.
- MAHLER, TAYLOR. 2019. Does at-issuehood predict projection? It's complicated! *North East Linguistic Society (NELS) 49*.245–55.
- MAHLER, TAYLOR. 2020. The social component of projection behavior of clausal complements. *Proceedings of the LSA 5*.777–91. DOI: 10.3765/plsa.v5i1.4703.
- MAHLER, TAYLOR; MARIE-CATHERINE DE MARNEFFE; and CATHERINE LAI. 2020. The prosody of presupposition projection in naturally-occurring utterances. *Proceedings of Sinn und Bedeutung 24(2)*.20–37. DOI: 10.18148/sub/2020.v24i2.884.
- PAVLICK, ELLIE, and TOM KWIAKOWSKI. 2019. Inherent disagreements in human textual inferences. *Transactions of the Association for Computational Linguistics* 7.677–94. DOI: 10.1162/tacl_a_00293.
- PINKER, STEVEN. 1994. *The language instinct*. New York: HarperPerennial. [Originally published by William Morrow and Co., 1994.]
- POTTS, CHRISTOPHER. 2005. *The logic of conventional implicatures*. Oxford: Oxford University Press. DOI: 10.1093/acprof:oso/9780199273829.001.0001.
- QING, CIYANG; NOAH D. GOODMAN; and DANIEL LASSITER. 2016. A rational speech-act model of projective content. *Proceedings of the 38th annual meeting of the Cognitive Science Society (CogSci 2016)*, 1110–15. Online: <https://cogsci.mindmodeling.org/2016/papers/0200/>.
- R CORE TEAM. 2016. R: A language and environment for statistical computing. Vienna: R Foundation for Statistical Computing. Online: <https://www.r-project.org/>.
- RICHARDSON, SYLVIA, and PETER J. GREEN. 1997. On Bayesian analysis of mixtures with an unknown number of components. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 59.731–92. DOI: 10.1111/1467-9868.00095.
- ROMOLI, JACOPO. 2015. The presuppositions of soft triggers are obligatory scalar implicatures. *Journal of Semantics* 32(2).173–219. DOI: 10.1093/jos/fft017.
- ROORYCK, JOHAN. 2000. *Configurations of sentential complementation*. London: Routledge.
- ROSS, ALEXIS, and ELLIE PAVLICK. 2019. How well do NLI models capture verb veridicality? *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2230–40. DOI: 10.18653/v1/D19-1228.
- SCHLENKER, PHILIPPE. 2003a. A plea for monsters. *Linguistics and Philosophy* 26.29–120. DOI: 10.1023/A:1022225203544.
- SCHLENKER, PHILIPPE. 2003b. The lazy Frenchman's approach to the subjunctive: Speculations on reference to worlds and semantics defaults in the analysis of mood. *Romance languages and linguistic theory 2003: Selected papers from 'Going Romance' 2003, Nijmegen, 20–22 November*, ed. by Twan Geerts, Ivo van Ginneken, and Haire Jacobs, 269–309. Amsterdam: John Benjamins.
- SCHLENKER, PHILIPPE. 2010. Local contexts and local meanings. *Philosophical Studies* 151.115–42. DOI: 10.1007/s11098-010-9586-0.
- SCRUCCA, LUCA; MICHAEL FOP; T. BRENDAN MURPHY; and ADRIAN E. RAFTERY. 2016. mclust 5: Clustering, classification and density estimation using Gaussian finite mixture models. *The R Journal* 8(1).289–317. DOI: 10.32614/RJ-2016-021.

- SIMONS, MANDY. 2007. Observations on embedding verbs, evidentiality, and presupposition. *Lingua* 117(6).1034–56. DOI: 10.1016/j.lingua.2006.05.006.
- SIMONS, MANDY; DAVID BEAVER; CRAIGE ROBERTS; and JUDITH TONHAUSER. 2017. The best question: Explaining the projection behavior of factive verbs. *Discourse Processes* 54(3).187–206. DOI: 10.1080/0163853X.2016.1150660.
- SIMONS, MANDY; JUDITH TONHAUSER; DAVID BEAVER; and CRAIGE ROBERTS. 2010. What projects and why. *Proceedings of Semantics and Linguistic Theory (SALT)* 20.309–27. DOI: 10.3765/salt.v20i0.2584.
- SMITH, E. ALLYN, and KATHLEEN CURRIE HALL. 2011. Projection diversity: Experimental evidence. *Proceedings of ESSLLI 2011 Workshop on Projective Content*, 156–70.
- SORACE, ANTONELLA, and FRANK KELLER. 2005. Gradience in linguistic data. *Lingua* 115(11).1497–1524. DOI: 10.1016/j.lingua.2004.07.002.
- SPECTOR, BENJAMIN, and PAUL EGRE. 2015. A uniform semantics for embedding interrogatives: An answer, not necessarily the answer. *Synthese* 192.1729–84. DOI: 10.1007/s11229-015-0722-4.
- SPROUSE, JON. 2007. Continuous acceptability, categorical grammaticality, and experimental syntax. *Biolinguistics* 1.123–34. DOI: 10.5964/bioling.8597.
- SPROUSE, JON, and DIOGO ALMEIDA. 2013. The role of experimental syntax in an integrated cognitive science of language. *The Cambridge handbook of biolinguistics*, ed. by Cedric Boeckx and Kleantzes K. Grohmann, 181–202. Cambridge: Cambridge University Press. DOI: 10.1017/CBO9780511980435.013.
- STEVENS, JON SCOTT; MARIE-CATHERINE DE MARNEFFE; SHARI R. SPEER; and JUDITH TONHAUSER. 2017. Rational use of prosody predicts projection in manner adverb utterances. *Proceedings of the 39th annual meeting of the Cognitive Science Society (CogSci 2017)*, 1144–49. Online: <https://cogsci.mindmodeling.org/2017/papers/0221/>.
- SWANSON, ERIC. 2012. Propositional attitudes. *Semantics: An international handbook of natural language meaning*, vol. 2, ed. by Klaus von Heusinger, Claudia Maienborn, and Paul Portner, 1538–61. Berlin: De Gruyter Mouton. DOI: 10.1515/9783110255072.1538.
- TONHAUSER, JUDITH. 2016. Prosodic cues to speaker commitment. *Proceedings of Semantics and Linguistic Theory (SALT)* 26.934–60. DOI: 10.3765/salt.v26i0.3788.
- TONHAUSER, JUDITH. 2020. Projection variability in Paraguayan Guarani. *Natural Language & Linguistic Theory* 38.1263–1302. DOI: 10.1007/s11049-019-09462-x.
- TONHAUSER, JUDITH; DAVID BEAVER; and JUDITH DEGEN. 2018. How projective is projective content? Gradience in projectivity and at-issueness. *Journal of Semantics* 35(3).495–542. DOI: 10.1093/jos/ffy007.
- TONHAUSER, JUDITH; DAVID BEAVER; CRAIGE ROBERTS; and MANDY SIMONS. 2013. Toward a taxonomy of projective content. *Language* 89(1).66–109. DOI: 10.1353/lan.2013.0001.
- TONHAUSER, JUDITH; MARIE-CATHERINE DE MARNEFFE; SHARI R. SPEER; and JON STEVENS. 2019. On the information structure sensitivity of projective content. *Proceedings of Sinn und Bedeutung* 23(2).363–89. DOI: 10.18148/sub/2019.v23i2.619.
- VAN DER SANDT, ROB A. 1992. Presupposition projection as anaphora resolution. *Journal of Semantics* 9(4).333–77. DOI: 10.1093/jos/9.4.333.
- VAN GELDEREN, ELLY. 2004. *Grammaticalization as economy*. Amsterdam: John Benjamins.
- VEHTARI, AKI; ANDREW GELMAN; and JONAH GABRY. 2017. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing* 27(5).1413–32. DOI: 10.1007/s11222-016-9696-4.
- WATANABE, SUMIO. 2010. Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research* 11(116).3571–94. Online: <https://jmlr.org/papers/v11/watanabe10a.html>.
- WHITE, AARON STEVEN. 2021. On believing and hoping whether. *Semantics and Pragmatics* 14:6. DOI: 10.3765/sp.14.6.
- WHITE, AARON STEVEN, and KYLE RAWLINS. 2018. The role of veridicality and factivity in clause selection. *North East Linguistic Society (NELS)* 48.221–34.
- WHITE, AARON STEVEN; RACHEL RUDINGER; KYLE RAWLINS; and BENJAMIN VAN DURME. 2018. Lexicosyntactic inference in neural models. *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 4717–24. DOI: 10.18653/v1/D18-1501.

- WILLIAMS, ADINA; NIKITA NANGIA; and SAMUEL BOWMAN. 2018. A broad-coverage challenge corpus for sentence understanding through inference. *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human language technologies, vol. 1 (long papers)*, 1112–22. DOI: 10.18653/v1/N18-1101.
- WYSE, BRENDAN. 2010. Factive/non-factive predicate recognition within question generation systems. Milton Keynes: The Open University master's thesis. Online: <http://computing-reports.open.ac.uk/2009/TR2009-09.pdf>.
- XUE, JINGYANG, and EDGAR ONEA. 2011. Correlation between projective meaning and at-issue: An empirical study. *Proceedings of ESSLLI 2011 Workshop on Projective Content*, 171–84.
- ZANUTTINI, RAFFAELLA, and PAUL PORTNER. 2003. Exclamative clauses: At the syntax-semantics interface. *Language* 79(1).39–81. DOI: 10.1353/lan.2003.0105.

Degen
Department of Linguistics
Stanford University
450 Jane Stanford Way
Stanford, CA 94305
[jdegen@stanford.edu]

[Received 2 April 2020;
revision invited 16 August 2020;
revision received 13 December 2020;
revision invited 18 June 2021;
revision received 30 September 2021;
accepted pending revisions 8 December 2021;
revision received 10 January 2022;
accepted 14 January 2022]

Tonhauser
Department of Linguistics
University of Stuttgart
Keplerstr. 17
70174 Stuttgart, Germany
[judith.tonhauser@ling.uni-stuttgart.de]